

# Parallel Programming

Divide and Conquer, Cilk-style bounds

Lets look at a code example: sum the elements of a list

# Sequential Version

The first step of writing a **parallel** program is writing a **sequential** version:

- Helps validate our eventual parallel program is correct
  - by comparing results with the simpler, sequential version
- Evaluate the performance of our parallel program
  - we write parallel programs to improve performance!


# Adding Numbers - Sequentially

```
public static int sum(int[] input){  
    int sum = 0;  
    for(int i=0; i<input.length; i++){  
        sum += input[i];  
    }  
    return sum;  
}
```

# Parallelism idea

Idea: Have 4 threads simultaneously sum 1/4 of the array

- **Warning:** This is an inferior first approach




- Create 4 *thread objects*, each given a portion of the work
- Call **start()** on each thread object to actually *run* it in parallel
- *Wait* for threads to finish using **join()**
- Add together their 4 answers for the *final result*

# First attempt, part 1



```
class SumThread extends java.lang.Thread {  
  
    int lo; // arguments  
    int hi;  
    int[] arr;  
  
    int ans = 0; // result  
  
    SumThread(int[] a, int l, int h) {  
        lo=l; hi=h; arr=a;  
    }  
  
    public void run() { //override must have this type  
        for(int i=lo; i < hi; i++)  
            ans += arr[i];  
    }  
}
```



Because we must override a no-arguments/no-result `run`,  
we use fields to communicate across threads

# First attempt, continued (wrong)

```
class SumThread extends java.lang.Thread {
    int lo, int hi, int[] arr; // arguments
    int ans = 0; // result
    SumThread(int[] a, int l, int h) { ... }
    public void run() { ... } // override
}

int sum(int[] arr) { // can be a static method
    int len = arr.length;
    int ans = 0;
    SumThread[] ts = new SumThread[4];
    for(int i=0; i < 4; i++) // do parallel computations
        ts[i] = new SumThread(arr, i*len/4, (i+1)*len/4);

    for(int i=0; i < 4; i++) // combine results
        ans += ts[i].ans;
    return ans;
}
```

## Second attempt (still wrong)

```
class SumThread extends java.lang.Thread {
    int lo, int hi, int[] arr; // arguments
    int ans = 0; // result
    SumThread(int[] a, int l, int h) { ... }
    public void run() { ... } // override
}

int sum(int[] arr) { // can be a static method
    int len = arr.length;
    int ans = 0;
    SumThread[] ts = new SumThread[4];
    for(int i=0; i < 4; i++) { // do parallel computations
        ts[i] = new SumThread(arr, i*len/4, (i+1)*len/4);
        ts[i].start(); // start actually runs the thread in parallel
    }
    for(int i=0; i < 4; i++) // combine results
        ans += ts[i].ans;
    return ans;
}
```



# Third attempt (correct in spirit)

```
class SumThread extends java.lang.Thread {
    int lo, int hi, int[] arr; // arguments
    int ans = 0; // result
    SumThread(int[] a, int l, int h) { ... }
    public void run() { ... } // override
}

int sum(int[] arr) { // can be a static method
    int len = arr.length;
    int ans = 0;
    SumThread[] ts = new SumThread[4];
    for(int i=0; i < 4; i++) { // do parallel computations
        ts[i] = new SumThread(arr, i*len/4, (i+1)*len/4);
        ts[i].start();
    }
    for(int i=0; i < 4; i++) { // combine results
        ts[i].join(); // wait for helper to finish!
        ans += ts[i].ans;
    }
    return ans;
}
```

# Discussion

The **Thread** class defines various methods you could not implement on your own

- For example: **start**, which calls **run** in a new thread

The **join** method is valuable for coordinating this kind of computation

- Caller blocks until/unless the receiver is done executing (meaning the call to **run** finishes)
- Else we would have a **race condition** on **ts[i].ans**

This style of parallel programming is called fork/join

Java detail: code has 1 compile error because **join** may throw **java.lang.InterruptedException**

- In basic parallel code, should be fine to catch-and-exit

# Shared memory?

Fork-join programs (thankfully) do not require much focus on sharing memory among threads

But in languages like Java, there is memory being shared.

In our example:

- **lo**, **hi**, **arr** fields written by “main” thread, read by helper thread
- **ans** field written by helper thread, read by “main” thread

When using shared memory, you must avoid race conditions (we will see a more formal definition of data races, later)

# Issues with this approach (and some workarounds)

Several reasons why this is a poor parallel algorithm

## Reason 1: want code to be reusable and efficient across platforms

- “Forward-portable” as core count grows
- So at the very least, **parameterize by the number of threads**

```
int sum(int[] arr, int numTs) {  
    int ans = 0;  
    SumThread[] ts = new SumThread[numTs];  
    for(int i=0; i < numTs; i++) {  
        ts[i] = new SumThread(arr, (i*arr.length)/numTs,  
                               ((i+1)*arr.length)/numTs);  
        ts[i].start();  
    }  
    for(int i=0; i < numTs; i++) {  
        ts[i].join();  
        ans += ts[i].ans;  
    }  
    return ans;  
}
```

# Issues with this approach (and some workarounds)

## **Reason 2:** want to use (only) processors “available to you *now*”

- Not used by other programs or threads in your program
  - Maybe caller is also using parallelism
  - Available cores can change even while your threads run

```
// numThreads == numProcessors is bad
// if some are needed for other things
int sum(int[] arr, int numTs){
    ...
}
```

# Issues with this approach (and some workarounds)

**Reason 3:** Though unlikely for **sum**, in general subproblems may take significantly different amounts of time

Example: Apply method **f** to every array element, but maybe **f** is much slower for some data items, e.g.: is a large integer prime?


If we create 4 threads and all slow data is processed by 1 of them, we won't get nearly a 4x speedup

- Example of a **load imbalance**

# A Better Approach

The counterintuitive (?) solution to all these problems is to use lots of threads, far more than the number of processors

- But this will require changing our algorithm
- And for constant-factor reasons, abandoning Java's threads




1. Forward-portable: Lots of helpers each doing a small piece
2. Processors available: Hand out “work chunks” as you go
3. Load imbalance: No problem if slow thread scheduled early enough
  - Variation probably small anyway if pieces of work are small

# Divide and Conquer to the Rescue!

This is straightforward to implement using divide-and-conquer

- Parallelism for the recursive calls





# Divide and Conquer

Fundamental pattern in parallel programming, also called **recursive splitting**

```
Divide and Conquer:  
  if cannot divide:  
    return unitary solution (stop recursion)  
  divide problem in two  
  solve first (recursively)  
  solve second (recursively)  
  combine solutions  
  return result
```

# Sequential Version: Recursive Sum

```
public static int do_sum_rec(int[] xs, int l, int h) {  
    int size = h-1;  
    if (size == 1) /*check for termination criteria*/  
        return xs[l];  
  
    /* split array in half and call self recursively*/  
    int mid = size / 2;  
    int sum1 = do_sum_rec(xs, l, l + mid);  
    int sum2 = do_sum_rec(xs, l + mid, h);  
    return sum1 + sum2;  
}
```

# Parallel Recursive Sum (with Threads)

```
public class SumThread extends Thread {  
    int[] xs;  
    int h, l;  
    int result;  
  
    public SumThread(int[] xs, int l, int h){  
        super();  
        this.xs = xs;  
        this.h = h;  
        this.l = l;  
    }  
  
    public void run(){  
        /*Do computation and write to result*/  
        return;  
    }  
}
```

# Parallel Recursive Sum (with Threads)

```
public void run(){
    int size = h-1;
    if (size == 1) {
        result = xs[1];
        return;
    }
    int mid = size / 2;
    SumThread t1 = new SumThread(xs, 1, 1 + mid);
    SumThread t2 = new SumThread(xs, 1 + mid, h);

    t1.start();
    t1.join();

    t2.start();
    t2.join();

    result = t1.result + t2.result;
    return;
}
```

Is this OK?

# Parallel Recursive Sum (with Threads)

```
public void run(){
    int size = l-h;
    if (size == 1) {
        result = xs[l];
        return;
    }
    int mid = size / 2;
    SumThread t1 = new SumThread(xs, l, l + mid);
    SumThread t2 = new SumThread(xs, l + mid, h);

    t1.start();
    t2.start();

    t1.join();
    t2.join();

    result = t1.result + t2.result;
    return;
}
```

Remark: This doesn't compile because join() can throw exceptions. In reality we need a try-catch block here.

# Result

**Java.lang.OutOfMemoryError: unable to create new native thread**

# One thread per parallel task model

Java threads are actually quite heavyweight

Java threads are mapped to OS threads (in the Oracle and most real-world implementations)

In general: using one thread per (small tasks) is highly inefficient

# Divide-and-Conquer works – (really, we'll get there)

In theory, you can divide down to single elements, do all your result-combining in parallel and get optimal speedup

In practice, creating all those threads and communicating swamps the savings, so:

- Use a *sequential cutoff*, typically around 500-1000
  - Eliminates *almost all* the recursive thread creation (bottom levels of tree)
- Do not create two recursive threads; create one and do the other “yourself”
  - Cuts the number of threads created by another 2x



# Divide-and-conquer – with manual fixes (Pt. I)

```
public void run(){
    int size = h-1;
    if (size < SEQ_CUTOFF)
        for (int i=1; i<h; i++)
            result += xs[i];
    else {
        int mid = size / 2;
        SumThread t1 = new SumThread(xs, 1, 1 + mid);
        SumThread t2 = new SumThread(xs, 1 + mid, h);
        t1.start();
        t2.start();
        t1.join();
        t2.join();
        result = t1.result + t2.result;
    }
}
```

# Half the threads

```
// wasteful: don't  
SumThread t1 = ...  
SumThread t2 = ...  
t1.start();  
t2.start();  
t1.join();  
t2.join();  
result=t1.result+t2.result;
```

```
// better: do  
// order of next 4 lines  
// essential - why?  
t1.start();  
t2.run();  
t1.join();  
result=t1.result+t2.result;
```

If a *language* had built-in support for fork-join parallelism, we would expect this hand-optimization to be unnecessary

But the *library* we are using expects you to do it yourself (and the difference is surprisingly substantial)

Again, no difference in theory


# Divide-and-conquer really works – (but it's hard work)

The key is divide-and-conquer parallelizes the result-combining

- *If* you have enough processors, total time is height of the tree:  $O(\log n)$  (optimal, exponentially faster than sequential  $O(n)$ )
- Often relies on operations being associative (like +)

Will write all our parallel algorithms in this style

- But using special libraries engineered for this style
  - Takes care of scheduling the computation well




# Recap: One thread per task model

Java threads are actually quite heavyweight

Java threads are mapped to OS threads


In general: using one thread per (small tasks) is highly inefficient

# Alternative approach: schedule tasks on threads




How many threads would you use?

# Java's executor service: managing asynchronous tasks



# Java's executor service:managing asynchronous tasks



```
.submit(Callable<T> task) → Future<T>  
.submit(Runnable task)   → Future<?>
```

# Note: Callable vs Runnable

ExecutorService can handle “Runnable” or “Callable” tasks:

Interface Runnable:

→ void run()



Does not return result

Interface Callable<T>:

→ T call()



Returns result



# Using executor service: Hello World (task)

```
static class HelloTask implements Runnable {  
  
    String msg;  
  
    public HelloTask(String msg) {  
        this.msg = msg;  
    }  
  
    public void run() {  
        long id = Thread.currentThread().getId();  
        System.out.println(msg + " from thread:" + id);  
    }  
}
```

# Using executor service: Hello World (creating executor, submitting)

```
int ntasks = 1000;
ExecutorService exs = Executors.newFixedThreadPool(4);

for (int i=0; i<ntasks; i++) {
    HelloTask t = new HelloTask("Hello from task " + i);
    exs.submit(t);
}

exs.shutdown(); // initiate shutdown, does not wait, but can't submit more tasks
```

# Using executor service: Hello World (output)

...

```
Hello from task 803 from thread:8  
Hello from task 802 from thread:10  
Hello from task 807 from thread:8  
Hello from task 806 from thread:9  
Hello from task 805 from thread:11  
Hello from task 810 from thread:9  
Hello from task 809 from thread:8  
Hello from task 808 from thread:10  
Hello from task 813 from thread:8  
Hello from task 812 from thread:9  
Hello from task 811 from thread:11
```

...

# Recursive Sum with ExecutorService


```
public Integer call() throws Exception {  
    int size = h - 1;  
    if (size == 1)  
        return xs[1];  
  
    int mid = size / 2;  
    sumRecCall c1 = new sumRecCall(ex, xs, 1, 1 + mid);  
    sumRecCall c2 = new sumRecCall(ex, xs, 1 + mid, h);  
  
    Future<Integer> f1 = ex.submit(c1);  
    Future<Integer> f2 = ex.submit(c2);  
  
    return f1.get() + f2.get();  
}
```


# Simple! – But does this work?


If you execute the code, you will observe that it never returns (i.e., the computation is not completed)




# Why does this happen?

 `sum(0,100):`  
    `t1 = spawn sum(0,50)`  
    `t2 = spawn sum(50,100)`  
    `t1.wait(); t2.wait()`

 `sum(0,50):`  
    `t1 = spawn sum(0,25)`  
    `t2 = spawn sum(25,50)`  
    `t1.wait(); t2.wait()`

 `sum(50,100):`  
    `t1 = spawn sum(50,75)`  
    `t2 = spawn sum(75,100)`  
    `t1.wait(); t2.wait()`


 `sum(0,25):`  
    `t1 = spawn sum(0,12)`  
    `t2 = spawn sum(12,25)`  
    `t1.wait(); t2.wait()`

tasks will end up waiting  
eventually we will run out of threads

# Adding Numbers ExecutorService: another approach

Problem with the divide and conquer approach is that tasks create other tasks and work partitioning (splitting up work) is part of the task.

A possible approach is to decouple work partitioning from solving the problem. That is we split the array into chunks (how many?) and create a task per chunk. Then, we submit tasks into ExecutorService and combine results (e.g., sum). It can be tricky to do the initial partitioning of work and final summing in parallel.




# Task Parallel Programming [Cilk-style]

## Tasks:

- execute code
- spawn other tasks
- wait for results from other tasks

A graph is formed based on spawning tasks



The edges mean that Task B was created by Task A and that Task C was created by Task A



# fib() Function

$$fib(n) = \begin{cases} n & n < 2 \\ fib(n-1) + fib(n-2) & n \geq 2 \end{cases}$$


## Sequential Version

```
public class Fibonacci {  
    public static long fib(int n){  
        if (n < 2)  
            return n;  
        long x1 = fib(n-1);  
        long x2 = fib(n-2);  
        return x1 + x2;  
    }  
}
```

## Parallel Version

```
public class Fibonacci {  
    public static long fib(int n) {  
        if (n < 2)  
            return n;  
        spawn task for fib(n-1);  
        spawn task for fib(n-2);  
        wait for tasks to complete  
        return addition of task results  
    }  
}
```

# fib(4) task graph




The task graph is a directed acyclic graph (DAG)

# Task parallelism discussion

- Tasks can execute in parallel
  - but they don't have to
  - assignment of tasks to CPUs/cores is up to the scheduler
- Task graph is **dynamic**
  - unfolds as execution proceeds
- Intuition: wide task graph → more parallelism

# Task parallelism: performance model


- Task graph: tasks become available as computation progresses
- We can execute the graph on  $p$  processors  
Scheduler assign tasks to processors
- $T_p$ : execution time on  $p$  processors



# Task parallelism: performance model


## [some reminders]

- $T_p$ : execution time on  $p$  processors
- $T_1$ : **work** (total amount of work)
  - the sum of the time cost of all nodes in graph
  - as if we executed graph sequentially ( $p=1$ )
- $T_1 / T_p \rightarrow$  **speedup**



# Task parallelism: performance model (Bounds)

- **$T_\infty$ : span, critical path**
  - Time it takes on infinite processors
  - longest path from root to sink
- **$T_1 / T_\infty \rightarrow$  parallelism**
  - “wider” is better
- Lower Bounds:
  - $T_p \geq T_1 / P$
  - $T_p \geq T_\infty$



On this graph,  $T_\infty$  is 8


# Scheduling of task graphs

Scheduler is an algorithm for assigning **tasks** to **processors**

Note that:



$T_p$  depends on scheduler

$T_1 / P$  and  $T_\infty$  are fixed




What is  $T_2$  for this graph?

That is, we have 2 processors.



$T_2$  will be 5 with  
this scheduling  
(we have 5 time steps)



$T_2$  will be 4 with  
this scheduling  
(we have 4 time steps)

a bound on how fast you can get on  $p$  processors  
with a greedy scheduler:  $T_p \leq T_1 / P + T_\infty$



# Work stealing scheduler

First used in MIT's Cilk, now a standard method

Provably:  $T_p = T_1 / P + O(T_\infty)$

Empirically:  $T_p \approx T_1 / P + T_\infty$

**Guideline for parallel programs => "Scheduling Multithreaded Computations by Work Stealing", Blumfoe & Leiserson, MIT**

# Summary

Divide and conquer for parallel programming

Cilk-style task graphs, scheduling and bounds