

TORSTEN HÖEFLER

# Parallel Programming, Spring 2021, Lecture 16: Spinlocks, Deadlocks, Semaphores

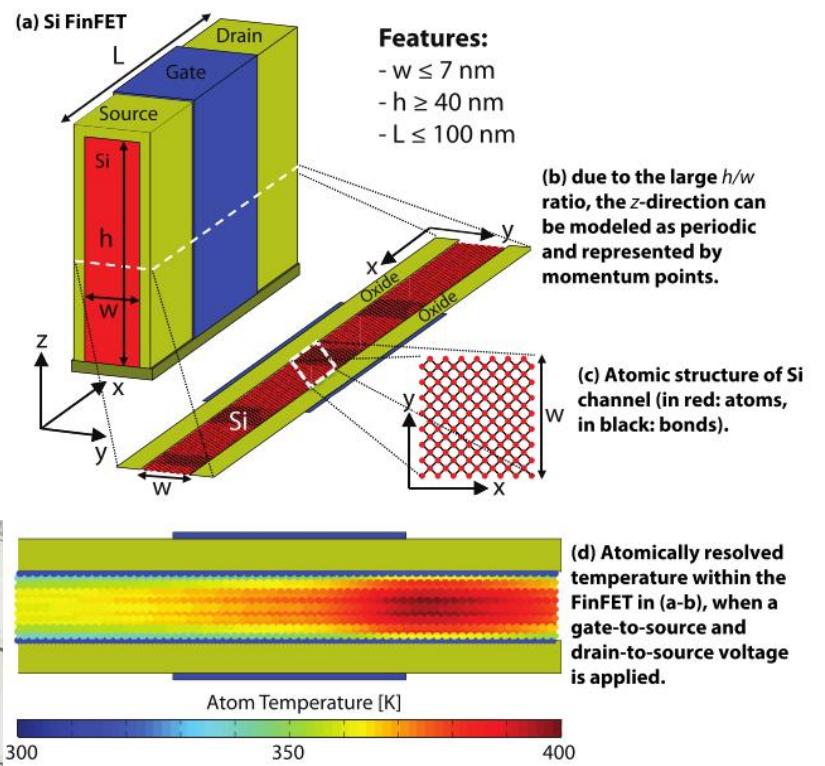
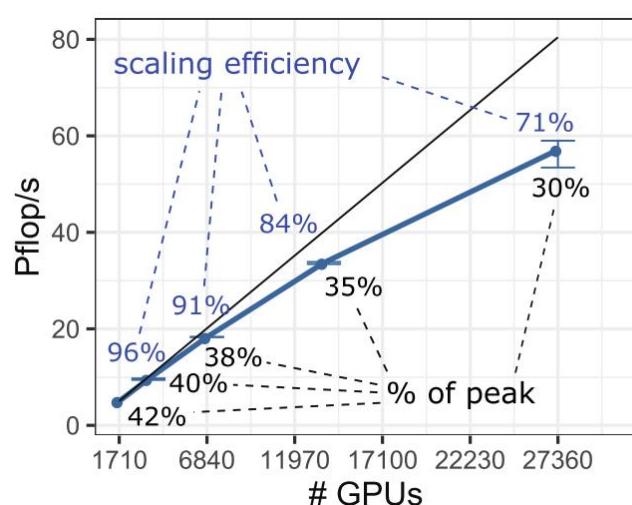


Figure 1: Simulation of self-heating effects in a 3-D Silicon FinFET with the OMEN code.

Some months later

Gordon Bell Award 2019



## ANNOUNCING THE WORLD'S FASTEST SUPERCOMPUTER FOR AI

20 Exaflops of AI

Powered by NVIDIA Grace CPU and  
Next Generation NVIDIA GPU

HPC and AI for Scientific and Commercial Apps

Advance Weather, Climate, and Material Science

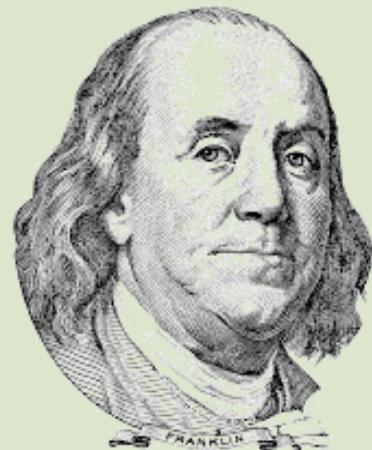


# Learning goals for today

- **So far:**
  - Bad interleavings and data races and why they happen
  - Memory ordering and how we formalize it to drive proofs
  - Implementation of locks using atomic or safe registers (Peterson lock)
- **Now:**
  - Multi-process locks (using SWMR registers)
  - Implementation of locks with Read-Modify-Write operations
  - Concurrency on a higher level: Deadlocks, Semaphores, Barriers
- **Learning goals:**
  - Understand pitfalls in very simple synchronization algorithms
  - This is very important to design correct parallel codes

## Recap last lecture by a short quiz

- Please participate in the zoom poll!



*“Tell me and I forget, teach me and I may remember,  
involve me and I learn.”*

# More concise than Decker: Peterson Lock (recap)

```
let P=1, Q=2; volatile boolean array flag[1..2] = [false, false];  
volatile integer victim = 1
```

## Process P (1)

loop

non-critical section

**flag[P] = true**

**victim = P**

**while(flag[Q] && victim == P);**

**critical section**

**flag[P] = false**

We both are interested

I am interested

but you go first

## Process Q (2)

loop

non-critical section

**flag[Q] = true**

**victim = Q**

**while(flag[P] && victim == Q);**

**critical section**

**flag[Q] = false**

We want to prove ...

that the Peterson Lock satisfies mutual exclusion  
and that it is starvation free

§ 1.<sup>7</sup> <sup>1</sup> Das kantonale Steueramt vollzieht das Erbschafts- und Schenkungssteuergesetz (ESchG) vom 28. September 1986<sup>4</sup>, soweit nachfolgend nichts Abweichendes geregelt ist.

<sup>2</sup> Die Finanzdirektion entscheidet über Rekurse gemäss §§ 61 Abs. 2 und 64 Abs. 2 ESchG<sup>4</sup>.

§ 2.<sup>6</sup> Es gelten sinngemäss:

- a. § 119 des Steuergesetzes<sup>2</sup> über den Ausstand,
- b. §§ 2–15, § 21 Abs. 1 und 2 sowie §§ 23–25 der Verordnung zum Steuergesetz<sup>3</sup>.

How?

Requires some notation first.

## Events and precedence

Threads produce a sequence of events (cf. actions)

P produces events  $p_0, p_1, \dots$

e.g.,  $p_1 = \text{"flag[P] = true"}$

programs usually consist of loops,  
therefore we might need to count  
occurrences

j-th occurrence of event i in thread P:  $p_i^j$

e.g.,  $p_5^3 = \text{"flag[P] = false"}$  in the third iteration

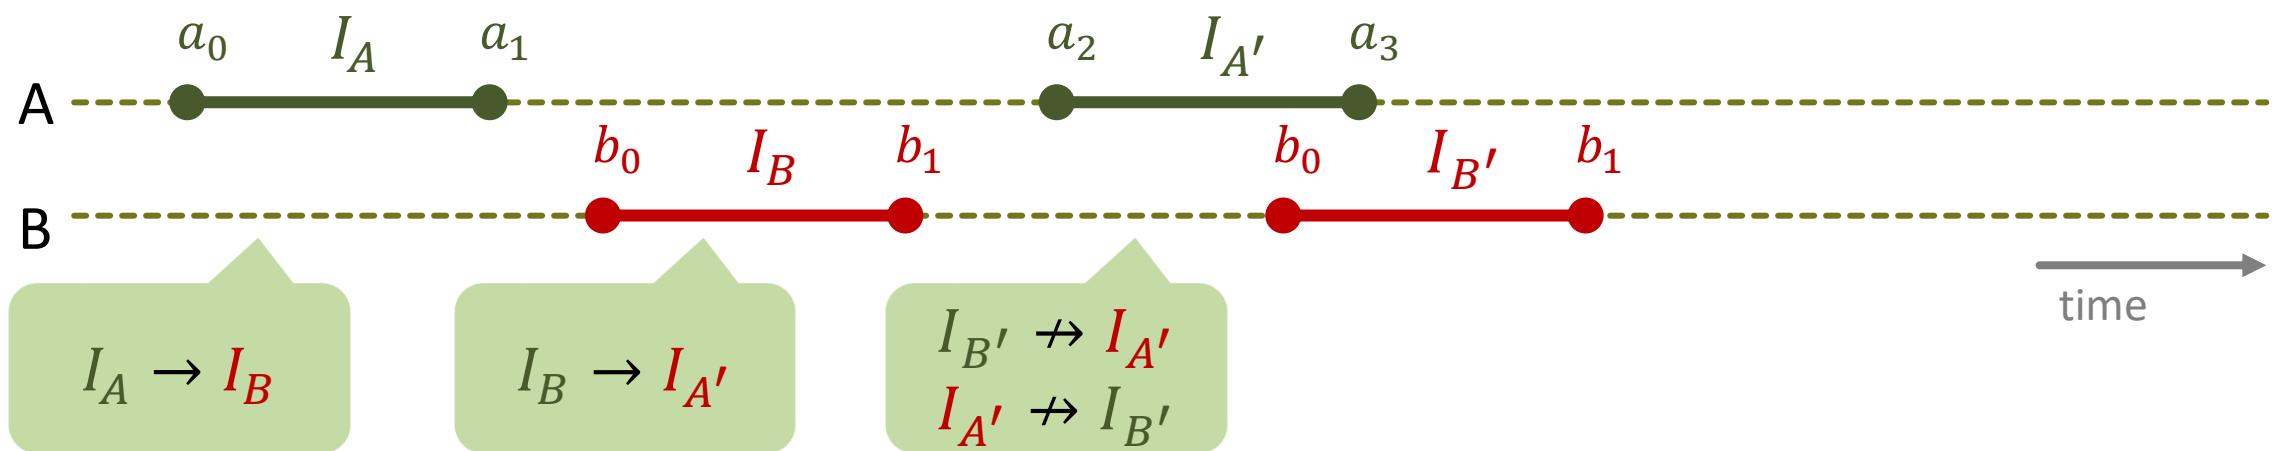
Precedence relation: we write  $a \rightarrow b$  when a occurs before b.

Note that the precedence relation " $\rightarrow$ " is a total order for events.

## Intervals

$(a_0, a_1)$ : interval of events  $a_0, a_1$  with  $a_0 \rightarrow a_1$

With  $I_A = (a_0, a_1)$  and  $I_B = (b_0, b_1)$  we write  $I_A \rightarrow I_B$  if  $a_1 \rightarrow b_0$



we say " $I_A$  precedes  $I_B$ " and " $I_{B'}$  and  $I_{A'}$  are concurrent"

## Atomic register

Register: basic memory object, can be shared or not

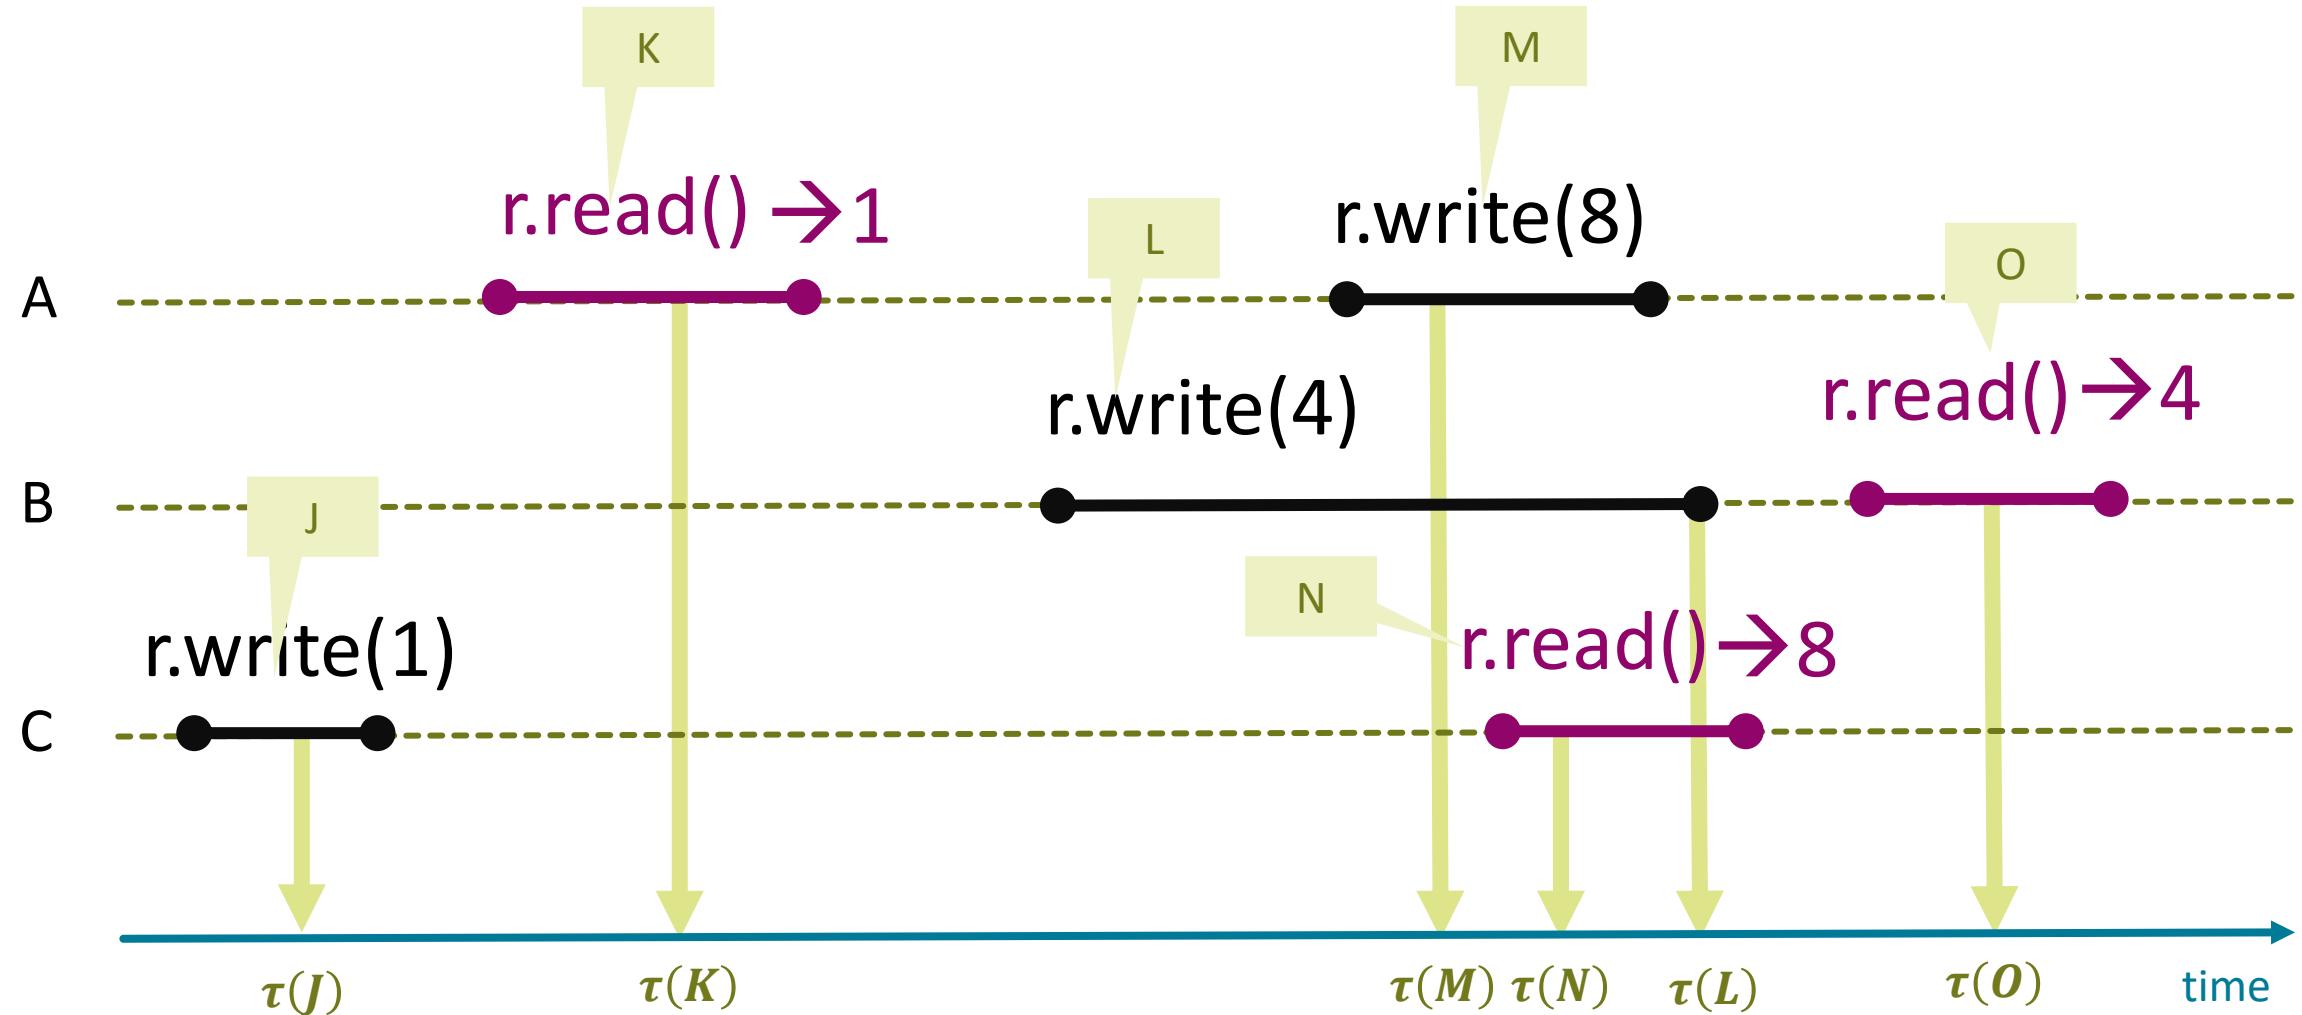
i.e., in this context register  $\neq$  register of a CPU

Register  $r$  : operations  $r.read()$  and  $r.write(v)$

Atomic Register:

- An invocation  $J$  of  $r.read$  or  $r.write$  takes effect at a single point  $\tau(J)$  in time
- $\tau(J)$  always lies between start and end of the operation  $J$
- Two operations  $J$  and  $K$  on the same register always have a different effect time  $\tau(J) \neq \tau(K)$
- An invocation  $J$  of  $r.read()$  returns the value  $v$  written by the invocation  $K$  of  $r.write(v)$  with closest preceding effect time  $\tau(K)$

## Example



## Atomic register

**Assumptions for atomic registers justify to treat operations on them as events taking place at a single point in time.**

**We will use this in the following proofs.**

**Note that even with atomic registers there can still be non-determinism of programs because nothing is said about the order of effect times for concurrent operations.**

## Proof: Mutual exclusion (Peterson)

By contradiction: assume concurrent  $CS_P$  and  $CS_Q$  [A]

Assume without loss of generality:

```
flag[P] = true
victim = P
while (flag[Q] && victim == P){}
CSP
flag[P] = false
```

$W_Q(\text{victim}=Q) \rightarrow W_P(\text{victim}=P)$  [B]

From the code:

$W_P(\text{flag}[P]=\text{true}) \rightarrow W_P(\text{victim} = P) \rightarrow R_P(\text{flag}[Q]) \rightarrow R_P(\text{victim}) \rightarrow CS_P$

"write of P"

A + C  $\Rightarrow$  must read false

B  $\Rightarrow$  must read P [C]

$W_Q(\text{flag}[Q]=\text{true}) \rightarrow W_Q(\text{victim} = Q) \rightarrow R_Q(\text{flag}[P]) \rightarrow R_Q(\text{victim}) \rightarrow CS_Q$

transitivity of " $\rightarrow$ "  
 $\Rightarrow$  must read true



"read of Q"

## Proof: Freedom from starvation

```
flag[P] = true
victim = P
while (flag[Q] && victim == P){}
CSP
flag[P] = false
```

**By (exhaustive) contradiction**

**Assume without loss of generality that P runs forever in its lock (while) loop, waiting until flag[Q]==false or victim != P.**

**Possibilities for Q:**

**stuck in nonCS**

⇒ flag[Q] = false and P can continue. Contradiction.

**repeatedly entering and leaving its CS**

⇒ sets victim to Q when entering.

Now victim cannot be changed ⇒ P can continue. Contradiction.

**stuck in its lock loop waiting until flag[P]==false or victim != Q.**

But victim == P and victim == Q cannot hold at the same time.

Contradiction.

# Peterson in Java

```
class PetersonLock
{
    volatile boolean flag[] = new boolean[2];
    volatile int victim;

    public void Acquire(int id)
    {
        flag[id] = true;
        victim = id;
        while (flag[1-id] && victim == id);
    }

    public void Release(int id)
    {
        flag[id] = false;
    }
}
```

Volatile reference to an array and not an array of volatile variables!  
This example may work in practice.  
However, for a correct program we need to use to use Java's **AtomicInteger** and **AtomicIntegerArray**.

# The Filter Lock

Extension of Peterson's lock to n processes

Every thread  $t$  knows its level in the filter  $level[t]$

In order to enter CS, a thread has to elevate all levels.

For each level, we use Peterson's mechanism to filter at most one thread, if other threads are at higher level.

For every level  $I$  there is one victim  $victim[I]$  that has to let others pass in case of conflicts.

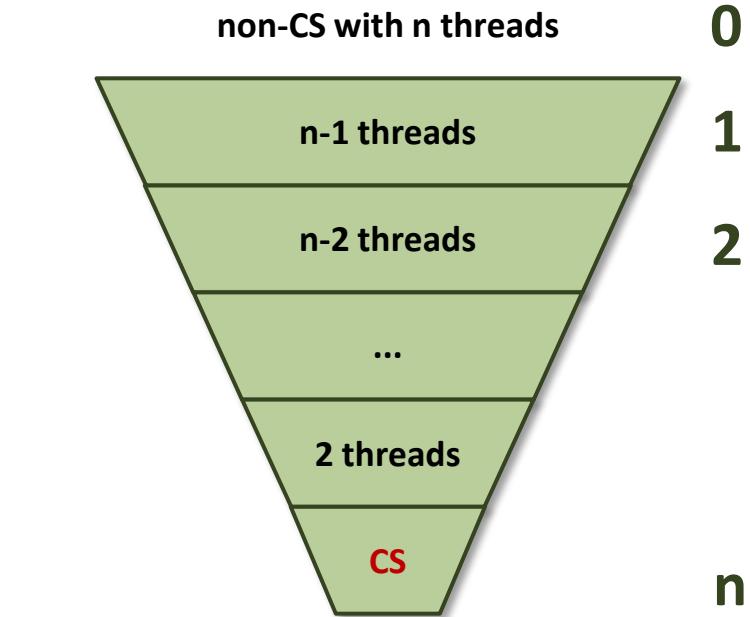


# The Filter Lock

```
int[] level(#threads), int[] victim(#threads)
```

```
lock(me) {  
    for (int i=1; i<n; ++i) {  
        level[me] = i;  
        victim[i] = me;  
        while ( $\exists k \neq me: level[k] \geq i \text{ && } victim[i] == me$ ) {};  
    }  
}  
  
unlock(me) {  
    level[me] = 0;  
}
```

Other threads  
are at same or  
higher level



And I have to wait

# FilterLock in Java

```
import java.util.concurrent.atomic.AtomicIntegerArray;
class FilterLock{
    AtomicIntegerArray level;
    AtomicIntegerArray victim;
    volatile int n;

    FilterLock(int n) {
        this.n = n;
        level = new AtomicIntegerArray(n);
        victim = new AtomicIntegerArray(n);
    }
    ...
}
```

# FilterLock in Java

```
...
//  $\exists k \neq me: level[k] \geq i (lev)$ 
boolean Others(int me, int lev) {
    for (int k = 0; k < n; ++k)
        if (k != me && level.get(k) >= lev) return true;
    return false;
}
public void Acquire(int me) {
    for (int lev = 1; lev < n; ++lev) {
        level.set(me, lev);
        victim.set(lev, me);
        while(me == victim.get(lev) && Others(me,lev));
    }
}
public void Release(int me) {
    level.set(me, 0);
}
```

Again: I (as a thread) can make progress if  
(a) Another thread wants to enter my level or  
(b) No more threads are in front of me  
This works because there are at most n  
threads in the system.

# Fairness

Divide lock implementation (preprotocol) into two parts

- doorway interval  $D$ : finite number of steps
- waiting interval  $W$ : unbounded number of steps

A lock algorithm is first-come-first-served when for two processes A and B it holds that

If  $D_A^j \rightarrow D_B^k$  then  $CS_A^j \rightarrow CS_B^k$



# The Filter Lock

satisfies mutual exclusion

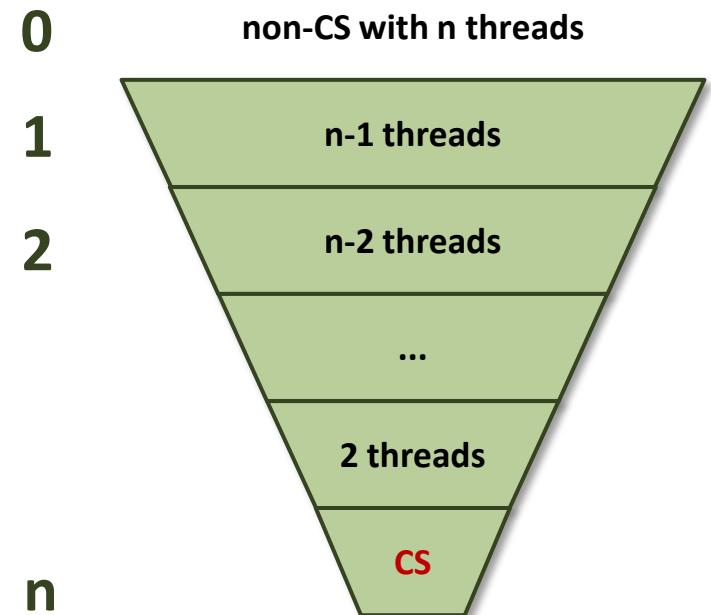
is deadlock free (how to prove?)

is starvation free (how to prove?)

but: is it also fair?

no: the filter lock is not first-come-first-serve

What else is bad about this lock?



## A small detour: Safe and Regular Registers

### Question:

- Is it possible to construct mutual exclusion with non-atomic registers?

Surprisingly: yes

- It is possible with registers fulfilling the weakest possible conditions that appear to be still useful in a concurrent setup.

## Weakening atomic registers – «the safe SWMR register»

Register  $r$ : basic memory object, can be shared or not,  
operations  $r.read()$  and  $r.write(v)$ .

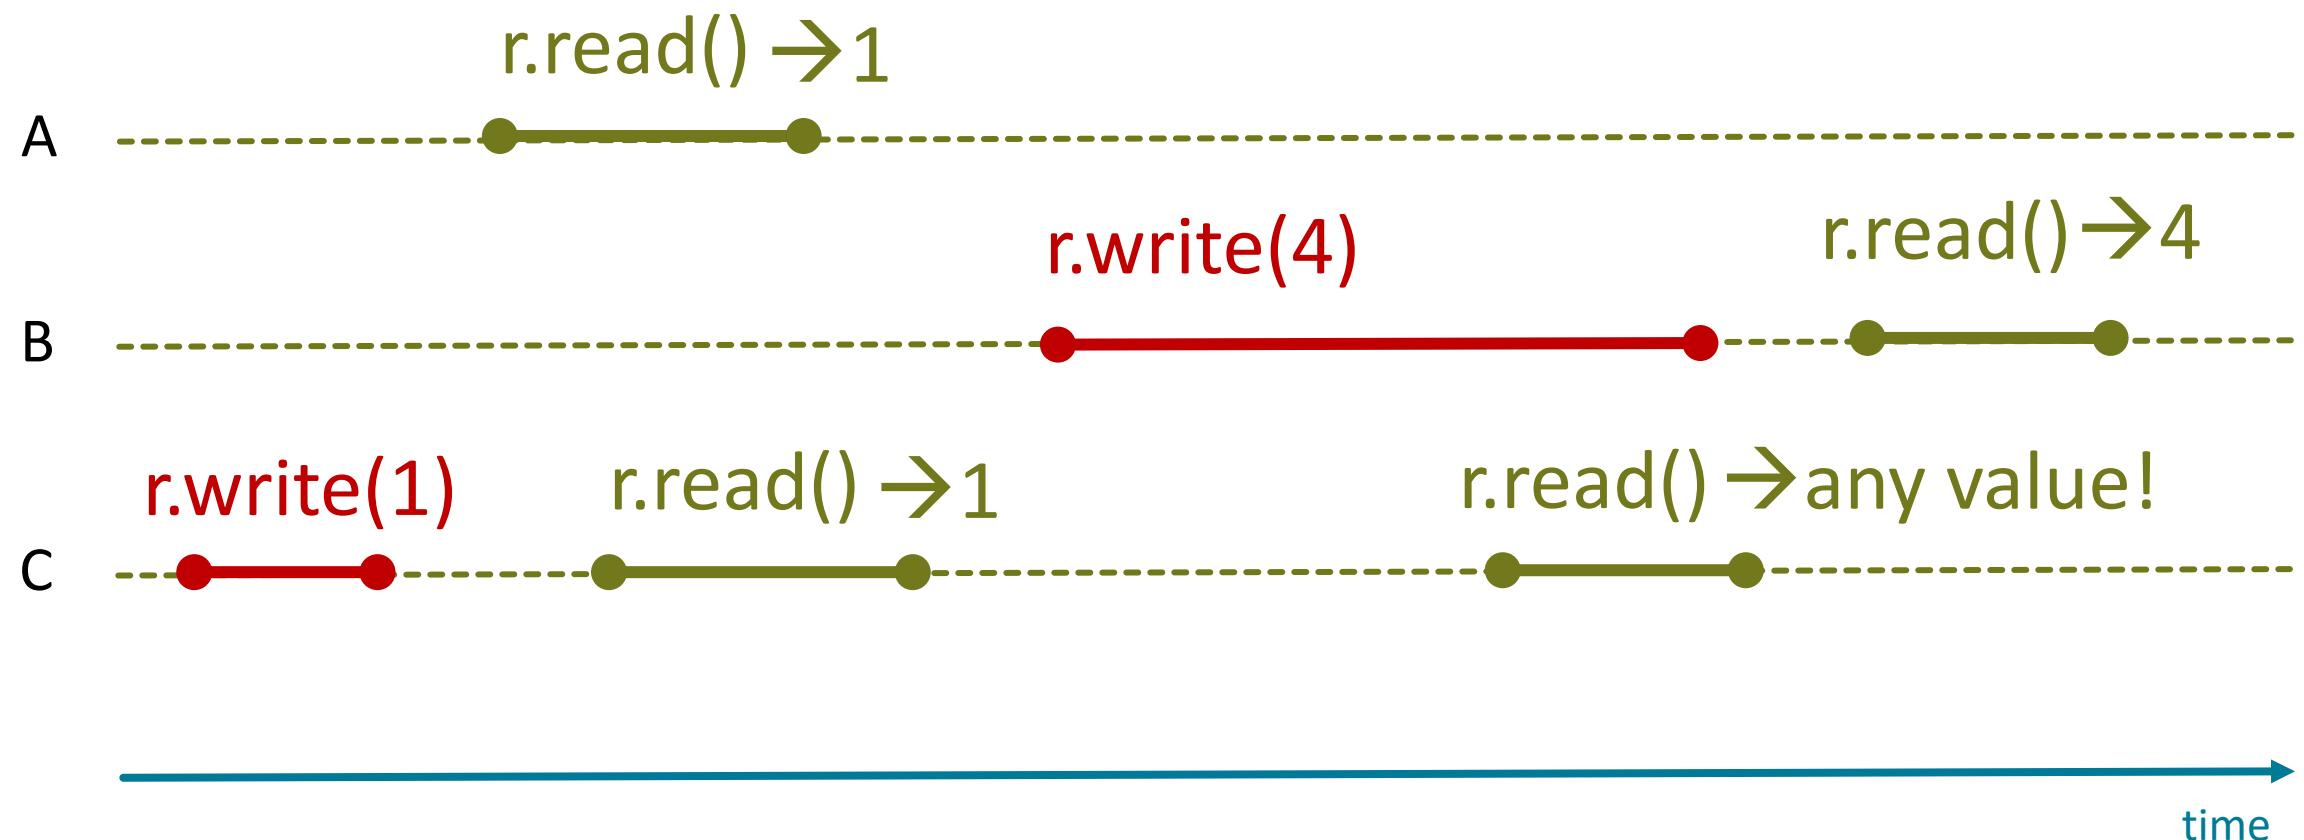
**SWMR (Single Writer Multiple Reader): only one concurrent write but multiple concurrent reads allowed.**

### *Safe* Register

- any read not concurrent with a write returns the current value of  $r$
- any read concurrent with a write can return *any value of the domain of r*  
if any read concurrent with writes can only return a value of one of the values (previous, new) then the register is called *regular*

The notion "safe" is historically motivated but actually misleading.

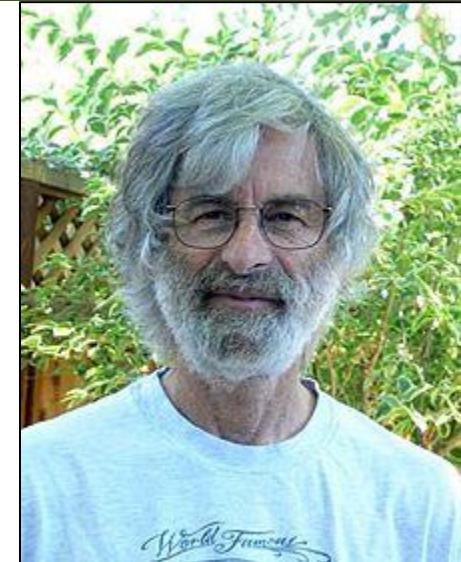
## Example



# Mutual Exclusion for n processes: Bakery Algorithm (1974)

A process is required to take a numbered ticket  
with value greater than all outstanding tickets

CS Entry: Wait until ticket number is lowest



Lamport, Turing award 2013



# Bakery algorithm (two processes, simplified)

```
volatile int np = 0, nq = 0
```

Process P

loop

non-critical section

$np = nq + 1$

while ( $nq \neq 0 \text{ && } nq < np$ );

critical section

$np = 0$

Q also wants access

and Q has an earlier ticket

Process Q

loop

non-critical section

$nq = np + 1$

while ( $np \neq 0 \text{ && } np \leq nq$ )

critical section

$nq = 0$

$np == nq$  can happen  
→ global ordering of processes

## Bakery algorithm (n processes)

```
integer array[0..n-1] label = [0, ..., 0]
boolean array[0..n-1] flag = [false, ..., false]
```

SWMR «ticket number»

SWMR «I want the lock»

**lock(me):**

```
flag[me] = true;
label[me] = max(label[0], ..., label[n-1]) + 1;
while ( $\exists k \neq me: flag[k] \&& (k, label[k]) <_l (me, label[me])$ ) {};
```

**unlock(me):**

```
flag[me] = false;
```

$$(k, l_k) <_l (j, l_j) \Leftrightarrow l_k < l_j \text{ or } (l_k = l_j \text{ and } k < j)$$

# Bakery Lock in Java

Nice lock! But which problem remains?

```
class BakeryLock
{
    AtomicIntegerArray flag; // there is no
    AtomicBooleanArray
    AtomicIntegerArray label;
    final int n;

    BakeryLock(int n) {
        this.n = n;
        flag = new AtomicIntegerArray(n);
        label = new AtomicIntegerArray(n);
    }

    int MaxLabel() {
        int max = label.get(0);
        for (int i = 1; i < n; ++i)
            max = Math.max(max, label.get(i));
        return max;
    }

    ...
}
```

```
boolean Conflict(int me) {
    for (int i = 0; i < n; ++i)
        if (i != me && flag.get(i) != 0) {
            int diff = label.get(i) - label.get(me);
            if (diff < 0 || diff == 0 && i < me)
                return true;
        }
    return false;
}

public void Acquire(int me) {
    flag.set(me, 1);
    label.set(me, MaxLabel() + 1);
    while(Conflict(me));
}

public void Release(int me) {
    flag.set(me, 0);
}
```

# In general

Shared memory locations (atomic registers) come with different access requirements

- Multi-Reader-Single-Writer (flag[], label[] in Bakery)
  - Multi-Reader-Multi-Writer (victim in Peterson)
- 
- Theorem 5.1 in [1]: *"If  $S$  is a [atomic] read/write system with at least two processes and  $S$  solves mutual exclusion with global progress [deadlock-freedom], then  $S$  must have at least as many variables as processes"*



INFORMATION AND COMPUTATION 107, 171–184 (1993)

[1]: Bounds on Shared Memory for Mutual Exclusion\*

JAMES E. BURNS

Georgia Institute of Technology, Atlanta, Georgia 30332

AND

NANCY A. LYNCH

Massachusetts Institute of Technology, Cambridge, Massachusetts 02139

The shared memory requirements of Dijkstra's mutual exclusion problem are examined. It is shown that  $n$  binary shared variables are necessary and sufficient to solve the problem of mutual exclusion with guaranteed global progress for  $n$  processes using only atomic reads and writes of shared variables for communication.



# We have constructed something ...

... that may not quite fulfil its purpose!

AND we cannot do better, can we?

- Is mutual exclusion really implemented like this?
  - NO! Why?
    - *space lower bound linear in the number of maximum threads!*
    - *without precautions (volatile variables) our assumptions on memory reordering does not hold. Memory barriers in hardware are expensive.*
    - *algorithms are not wait-free (more later)*
- But we proved that we cannot do better. What now?
  - Change (extend) the model with architecture engineering!
    - *modern multiprocessor architectures provide special instructions for atomically reading and writing at once!*



# Hardware Support for Parallelism

## Read-Modify-Write Operations

# Hardware support for atomic operations: Example (x86)

## CMPXCHG

Compares the value in the AL, AX, EAX, or RAX register with the value in a register or a memory location (first operand). If the two values are equal, the instruction copies the value in the second operand to the first operand and sets the ZF flag in the rFLAGS register to 1. Otherwise, it copies the value in the first operand to the AL, AX, EAX, or RAX register and clears the ZF flag to 0.

The OF, SF, AF, PF, and CF flags are set to reflect the results of the compare.

When the first operand is a memory operand, the second operand must also be a memory operand.

The forms of the first and second operands depend on the LOCK prefix.

The forms of the first and second operands depend on the LOCK prefix.

## Mnemonic

CMPXCHG

CMPXCHG

CMPXCHG

CMPXCHG reg/mem64, reg64

0F B1 /r

## Related Instructions

CMPXCHG8B, CMPXCHG16B

## Compare and Exchange

**CMPXCHG mem, reg**  
«compares the value in Register A with the value in a memory location. If the two values are equal, the instruction copies the value in the second operand to the first operand and sets the ZF flag in the flag registers to 1. Otherwise it copies the value in the first operand to the A register and clears ZF flag to 0»

Compare the value in a register or memory location with the value in a register or memory location. If equal, copy the second operand to the first operand. Otherwise, copy the first operand to RAX.

## 1.2.5 Lock Prefix

The LOCK prefix causes certain kinds of memory read-modify-write instructions to occur atomically. The mechanism for doing so is implementation-dependent (for example, the mechanism may involve

8

Instruction Formats

«The LOCK prefix causes certain kinds of memory read-modify-write instructions to occur atomically»

24594—Rev. 3.14—September 2007

AMD

AMD64 Technology

bus signaling or packet messaging between the processor and a memory controller). The prefix is intended to give the processor exclusive use of shared memory in a multiprocessor system.

The LOCK prefix can only be used with forms of the following instructions that write a memory operand: ADC, ADD, AND, BTC, BTR, BTS, CMPXCHG, CMPXCHG8B, CMPXCHG16B, DEC, INC, NEG, NOT, OR, SBB, SUB, XADD, XCHG, and XOR. An invalid-opcode exception occurs if the LOCK prefix is used with any other instruction.

From the AMD64 Architecture Programmer's Manual

# Hardware support for atomic operations: Example (ARM)

## LDREX

31	28	27	26	25	24	23	22	21	20	19	16	15	12	11	8	7	6	5	4	3	0
cond	0	0	0	1	1	0	0	1	Rn	Rd	SBO	1	0	0	1	SBO					

LDREX (Load Register Exclusive) loads a register from memory, and:

- if the address has the Shared memory attribute, mark the physical address as exclusive access for the executing processor.
- causes the executing processor to release the shared memory attribute.

**LDREX <rd>, <rn>**

«Loads a register from memory and if the address has the shared memory attribute, mark the physical address as exclusive access for the executing processor in a shared monitor»

## Syntax

LDREX{<cond>} <Rd>, [<Rn>]

where:

<cond> Is the condition code for the conditional operation.

<Rd> Specifies the register to be loaded.

<Rn> Specifies the register containing the address.

## Architecture version

Version 6 and above.

## STREX

31	28	27	26	25	24	23	22	21	20	19	16	15	12	11	8	7	6	5	4	3	0
cond	0	0	0	1	1	0	0	0	Rn	Rd	SBO	1	0	0	1	SBO	1	0	0	1	Rm

STREX (Store Register Exclusive) performs a conditional store to memory. The store only occurs if the executing processor has exclusive access to the memory addressed.

## Syntax

STREX{<cond>} <Rd>, <Rm> [<Rn>]

where:

<cond>

<Rd>

<Rm>

<Rn>

**STREX <rd>, <rm>, <rn>**

«performs a conditional store to memory. The store only occurs if the executing processor has exclusive access to the memory addressed»

if the operation updates memory  
if the operation fails to update memory.

are defined in The  
on is used.

returned is:

From the ARM Architecture Reference Manual

# Hardware support for atomic operations

## Typical instructions

### Test-And-Set (TAS)

Example TSL register, flag (Motorola 68000)

### Compare-And-Swap (CAS)

Example: LOCK CMPXCHG (Intel x86)

Example: CASA (Sparc)

### Load Linked / Store Conditional

Example LDREX/STREX (ARM)

Example LL / SC (MIPS, POWER, RISC V)

Atomic instructions are typically much slower than simple read & write operations [1]!

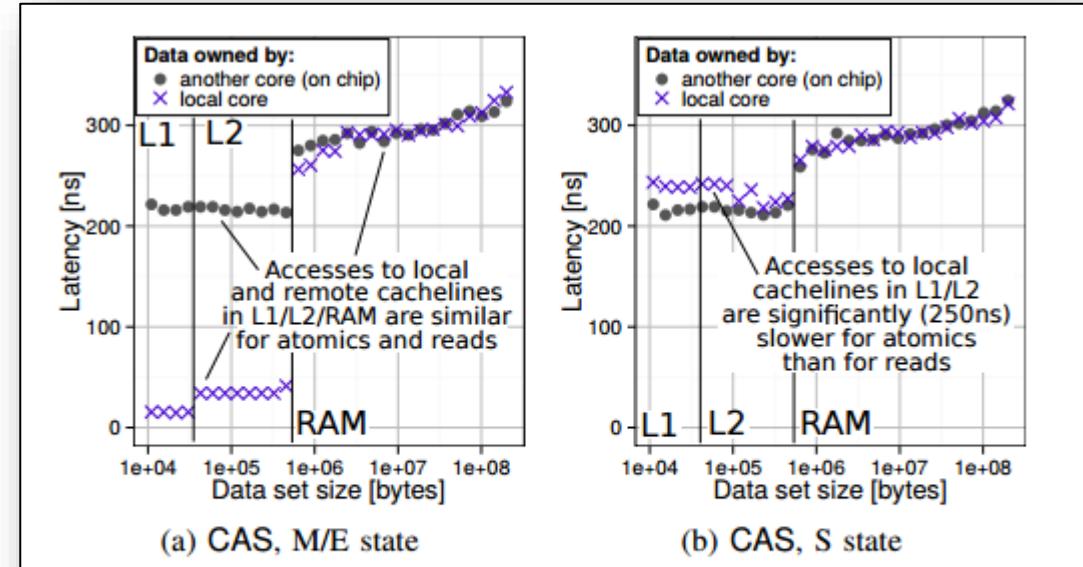


Fig. 6: The comparison of the latency of CAS on Xeon Phi. The requesting core accesses its own cache lines (local) and cache lines of a different core from the same chip (on chip).

## Semantics of TAS and CAS

```
boolean TAS(memref s)
    if (mem[s] == 0) {
        mem[s] = 1;
        return true;
    } else
        return false;
```

atomic

```
int CAS (memref a, int old, int new)
    oldval = mem[a];
    if (old == oldval)
        mem[a] = new;
    return oldval;
```

atomic

- are **Read-Modify-Write** («atomic») operations
- enable implementation of a mutex with **O(1)** space  
(in contrast to Filter lock, Bakery lock etc.)
- are needed for lock-free programming (later in this course)

# Implementation of a spinlock using simple atomic operations

## Test and Set (TAS)

### Init (lock)

```
lock = 0;
```

## Compare and Swap (CAS)

### Init (lock)

```
lock = 0;
```

### Acquire (lock)

```
while !TAS(lock); // wait
```

### Acquire (lock)

```
while (CAS(lock, 0, 1) != 0);
```

### Release (lock)

```
lock = 0;
```

### Release (lock)

```
CAS(lock, 1, 0);
```

ignore result

# Read-Modify-Write in Java

Let's try it.

Need support for atomic operations on a high level.

Available in Java (from JDK 5) with class

**java.util.concurrent.atomic.AtomicBoolean**

Operations

**boolean set();**

**boolean get();**

**boolean compareAndSet(boolean expect, boolean update);**

**boolean getAndSet(boolean newValue);**

atomically set to value **update** iff current value is **expect**. Return true on success.

sets **newValue** and returns previous value.

## How does this work? (for experts)

- The JVM bytecode does not offer atomic operations like CAS.  
[It does, however, support monitors via instructions monitorenter, monitorexit, we will understand this later]
- But there is a (yet undocumented) class **sun.misc.Unsafe** offering direct mappings from java to underlying machine / OS.
- Direct mapping to hardware is not guaranteed – operations on AtomicBoolean are not guaranteed lock-free

# Example: java.util.concurrent.atomic.AtomicInteger (for experts)

```
35
36 package java.util.concurrent.atomic;
37 import sun.misc.Unsafe;
```

---

...

Atomically sets the value to the given updated value if the current value == the expected value.

**Parameters:**

expect the expected value  
update the new value

**Returns:**

true if successful. False return indicates that the actual value was not equal to the expected value.

```
133
134     public final boolean compareAndSet(int expect, int update) {
135         return unsafe.compareAndSwapInt(this, valueOffset, expect, update);
136     }
```

(source: [grepcode.com](http://grepcode.com))

# TASLock in Java

```
public class TASLock implements Lock {  
    AtomicBoolean state = new AtomicBoolean(false);  
  
    public void lock() {  
        while(state.getAndSet(true)) {}  
    }  
  
    public void unlock() {  
        state.set(false);  
    }  
    ...  
}
```

**Spinlock:**

**Try to get the lock.**

**Keep trying until the lock is acquired (return value is false).**

**unlock**

**release the lock (set to false)**



# Simple TAS Spin Lock – Measurement Results

## TAS

n = 1, elapsed= 224, normalized= 224

n = 2, elapsed= 719, normalized= 359

n = 3, elapsed= 1914, normalized= 638

n = 4, elapsed= 3373, normalized= 843

n = 5, elapsed= 4330, normalized= 866

n = 6, elapsed= 6075, normalized= 1012

n = 7, elapsed= 8089, normalized= 1155

n = 8, elapsed= 10369, normalized= 1296

n = 16, elapsed= 41051, normalized= 2565

n = 32, elapsed= 156207, normalized= 4881

n = 64, elapsed= 619197, normalized= 9674

```
public class TASLock implements Lock {  
    AtomicBoolean state = new AtomicBoolean(false);  
  
    public void lock() {  
        while(state.getAndSet(true)) {}  
    }  
  
    public void unlock() {  
        state.set(false);  
    }  
}
```

- run n threads
- each thread acquires and releases the TASLock a million times
- repeat scenario ten times and add up runtime
- record time per thread

Intel core i7@3.4 GHz, 4 cores + HT

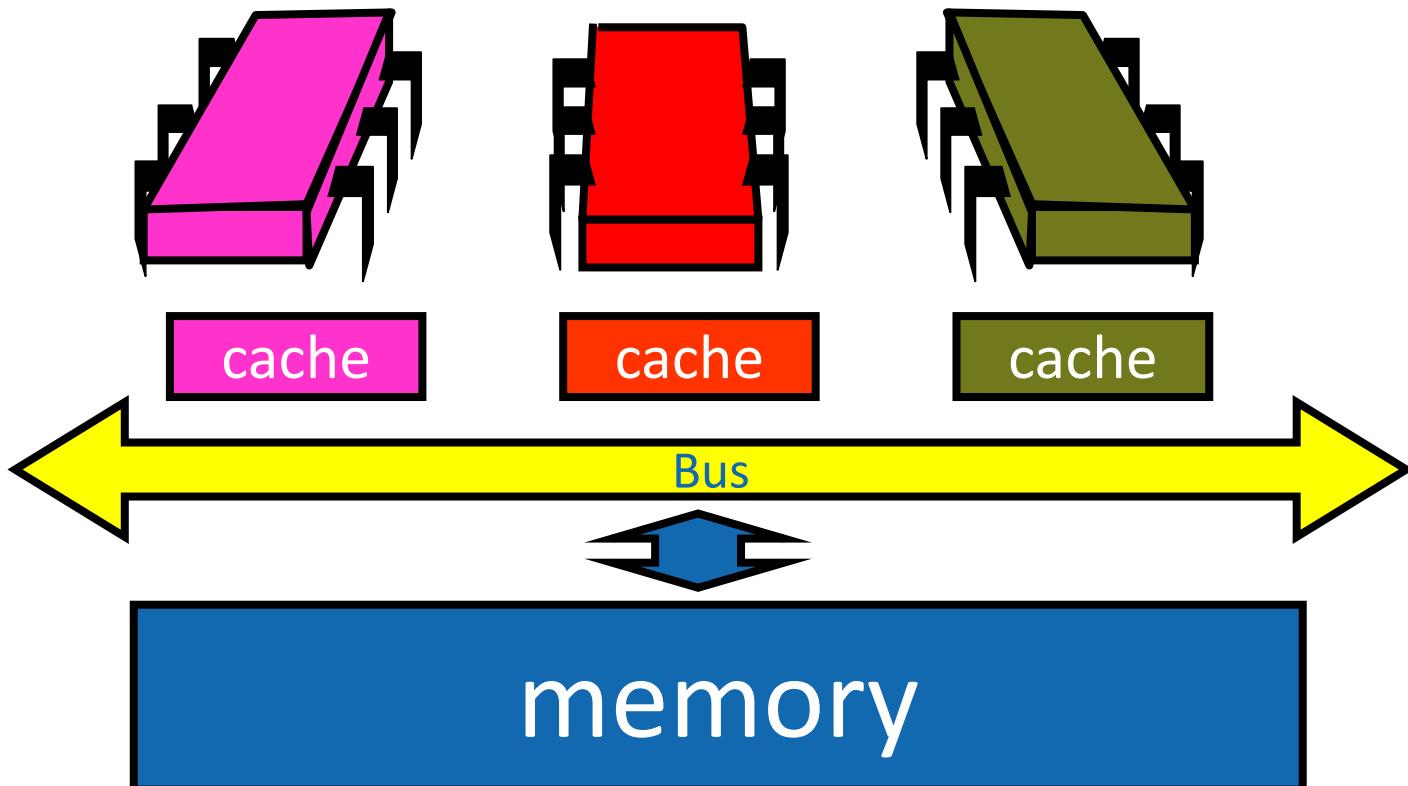
# Why?

**sequential bottleneck**

**contention:** threads fight  
for the bus during call of  
`getAndSet()`

cache coherency protocol  
invalidates cached copies  
of the lock variable on  
other processors

```
public class TASLock implements Lock {  
    AtomicBoolean state = new AtomicBoolean(false);  
  
    public void lock() {  
        while(state.getAndSet(true)) {}  
    }  
  
    public void unlock() {  
        state.set(false);  
    }  
}
```

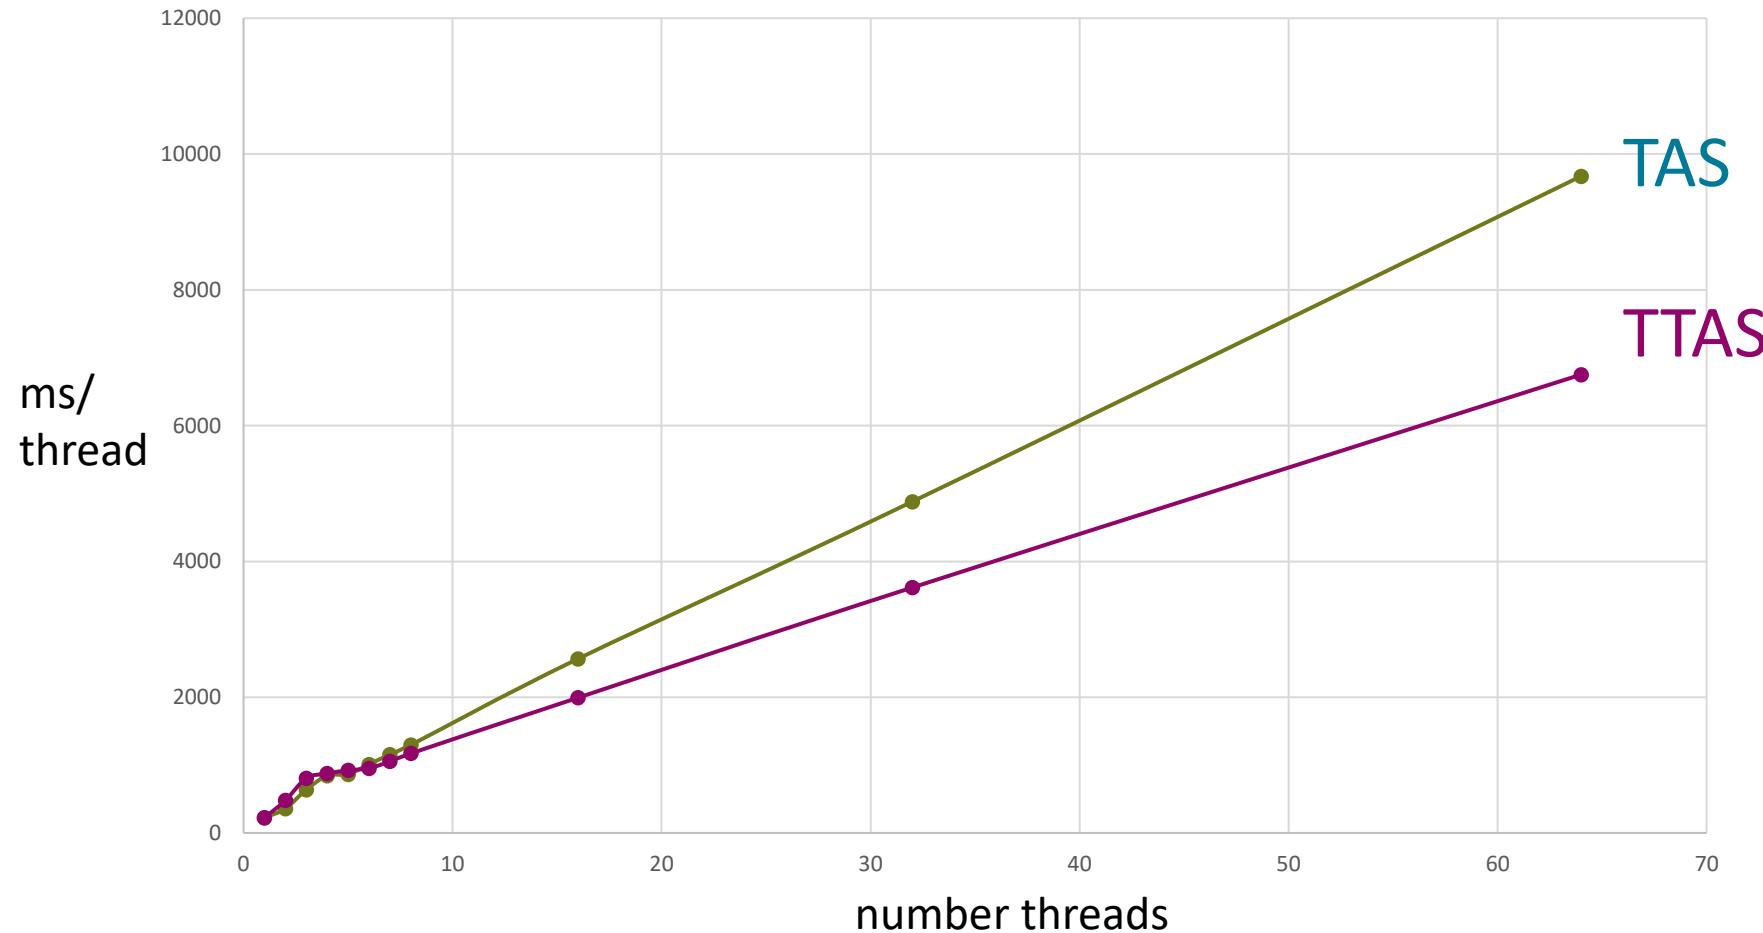


## Test-and-Test-and-Set (TATAS) Lock

```
public void lock()
{
    do
        while(state.get()) {}
    while (!state.compareAndSet(false, true));
}

public void unlock()
{
    state.set(false);
}
```

# Measurement



note that this varies strongly between machines and JVM implementations and even between runs.  
Take it as a qualitative statement

# TATAS does not generalize

## ■ Example: Double-Checked Locking

**Double-Checked Locking**

An Optimization Pattern for Efficiently Initializing and Accessing Thread-safe Objects

Douglas C. Schmidt  
schmidt@cs.wustl.edu  
Dept. of Computer Science  
Wash. U., St. Louis

Tim Harrison  
harrison@cs.wustl.edu  
Dept. of Computer Science  
Wash. U., St. Louis

This paper appeared in a chapter in the book "Pattern Languages of Program Design 3" ISBN, edited by Robert Martin, Frank Buschmann, and Dirk Riehle published by Addison-Wesley, 1997.

**Abstract**

*This paper shows how the canonical implementation [1] of the Singleton pattern does not work correctly in the presence of preemptive multi-tasking or true parallelism. To solve this problem, we present the Double-Checked Locking optimization pattern. This pattern is useful for reducing contention and synchronization overhead whenever "critical sections" of code should be executed just once. In addition, Double-Checked Locking illustrates how changes in underlying forces (i.e., adding multi-threading and parallelism to the common Singleton use-case) can impact the form and content of patterns used to develop concurrent software.*

```
class Singleton
{
public:
    static Singleton *instance (void)
    {
        if (instance_ == 0)
            // Critical section.
            instance_ = new Singleton;
        return instance_;
    }
}
```

double-checked locking

About 830,000 results (0.27 seconds)

**Double-checked locking - Wikipedia, the free encyclopedia**  
[en.wikipedia.org/wiki/Double-checked\\_locking](https://en.wikipedia.org/wiki/Double-checked_locking)  
In software engineering, **double-checked locking** (also known as "**double-checked locking** optimization") is a software design pattern used to reduce the ...  
Usage in Java - Usage in Microsoft Visual C++ - Usage in Microsoft .NET ...

**The "Double-Checked Locking is Broken" Declaration**  
[www.cs.umd.edu/~pugh/java/.../DoubleCheckedLocking.html](http://www.cs.umd.edu/~pugh/java/.../DoubleCheckedLocking.html)  
Details on the reasons - some very subtle - why **double-checked locking** cannot be relied upon to be safe. Signed by a number of experts, including Sun ...

**Double-checked locking and the Singleton pattern**  
[www.ibm.com/developerworks/java/library/j-dcli/index.html](http://www.ibm.com/developerworks/java/library/j-dcli/index.html)  
1 May 2002 – **Double-checked locking** is one such idiom in the Java programming language that should never be used. In this article, Peter Haggar ...

**Double-checked locking: Clever, but broken - JavaWorld**  
[www.javaworld.com > Java Development Tools](http://www.javaworld.com > Java Development Tools)  
9 Feb 2001 – Many Java programmers are familiar with the **double-checked locking** idiom, which allows you to perform lazy initialization with reduced ...

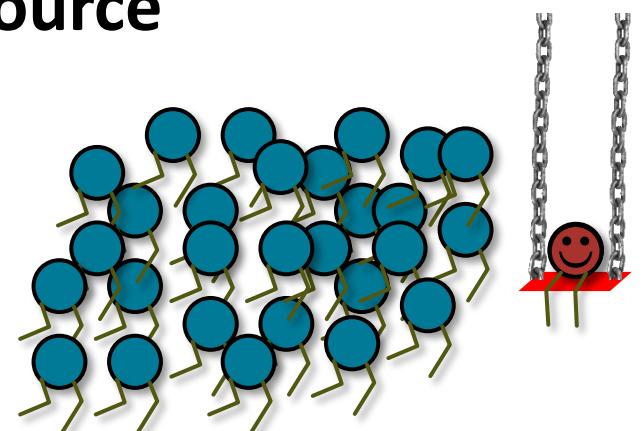
[PDF] **Double-Checked Locking An Optimization Pattern for Efficiently ...**  
[sunsite.icm.edu.pl/packages/ace/ACE/PDF/DC-Locking.pdf](http://sunsite.icm.edu.pl/packages/ace/ACE/PDF/DC-Locking.pdf)  
File Format: PDF/Adobe Acrobat - Quick View  
by DC Schmidt - Cited by 14 - Related articles  
solve this problem, we present the **Double-Checked Lock** ing optimization ...  
**Double-Checked Locking** illustrates how changes in under- lying forces (i.e. ...

Problem: Memory ordering leads to race-conditions!

# TATAS with backoff

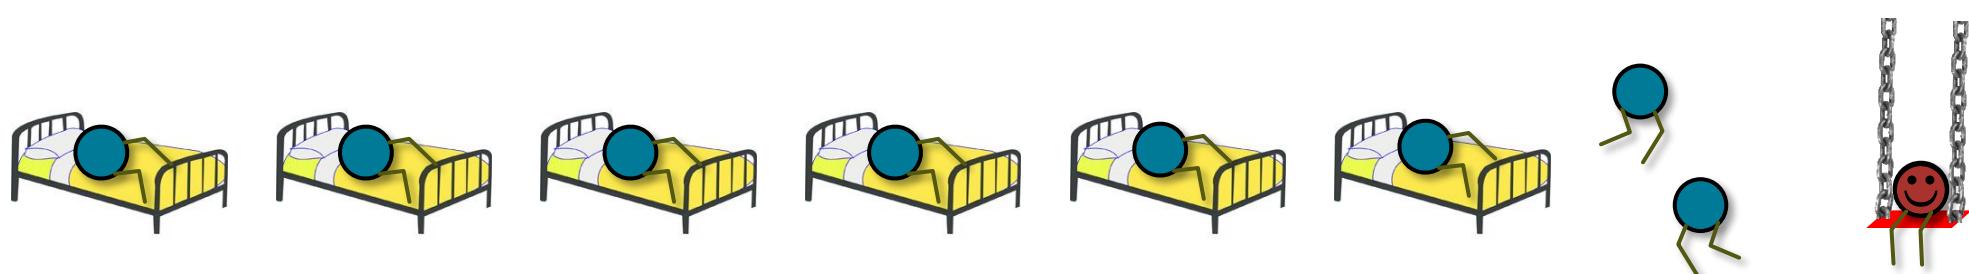
## Observation

- (too) many threads fight for access to the same resource
- slows down progress globally and locally



## Solution

- threads go to sleep with random duration
- increase expected duration each time the resource is not free



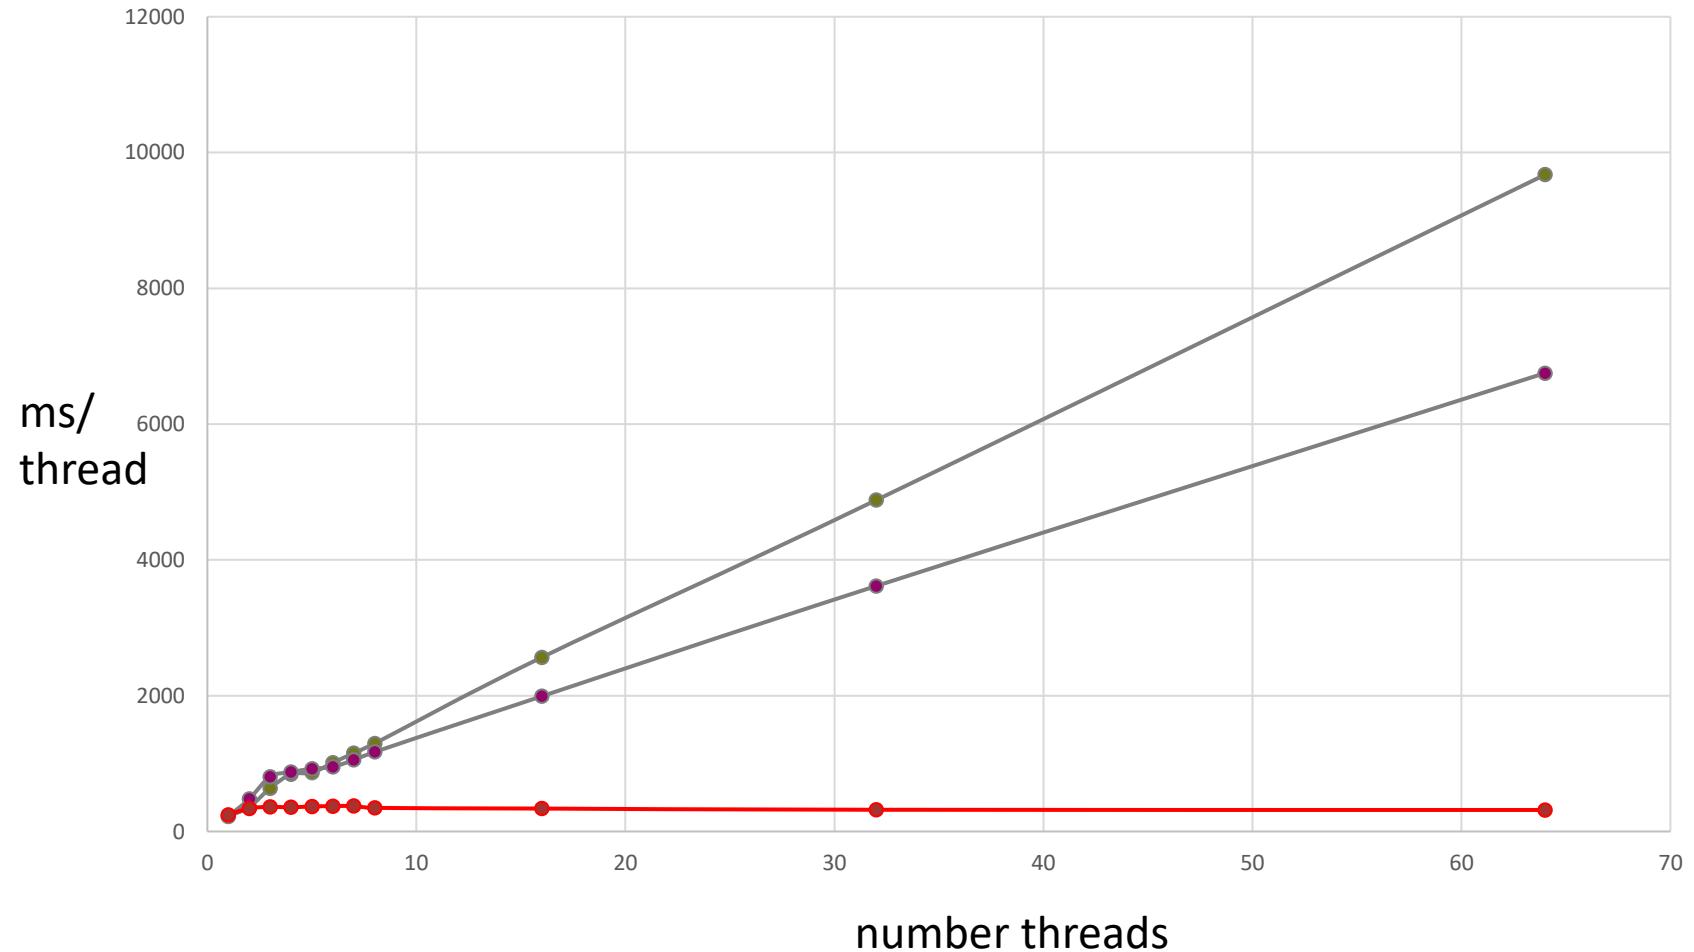
# Lock with Backoff

```
public void lock() {  
    Backoff backoff = null;  
    while (true) {  
        while (state.get()) {};           // spin reading only (TTAS)  
        if (!state.getAndSet(true))      // try to acquire, returns previous val  
            return;  
        else { // backoff on failure  
            try {  
                if (backoff == null) // allocation only on demand  
                    backoff = new Backoff(MIN_DELAY, MAX_DELAY);  
                backoff.backoff();  
            } catch (InterruptedException ex) {}  
        }  
    }  
}
```

# exponential backoff

```
class Backoff
{
    ...
    public void backoff() throws InterruptedException {
        int delay = random.nextInt(limit);
        if (limit < maxDelay) { // double limit if less than max
            limit = 2 * limit;
        }
        Thread.sleep(delay);
    }
}
```

# Measurement



TAS

TTAS

BackoffLock

yeah!

# Deadlock

## Deadlocks – Motivation

Consider a method to transfer money between bank accounts

```
class BankAccount {  
    ...  
    synchronized void withdraw(int amount) {...}  
    synchronized void deposit(int amount) {...}  
  
    synchronized void transferTo(int amount, BankAccount a) {  
        this.withdraw(amount);  
        a.deposit(amount);  
    }  
}
```

Thread aquires second lock in a.deposit.  
Can this become a problem?

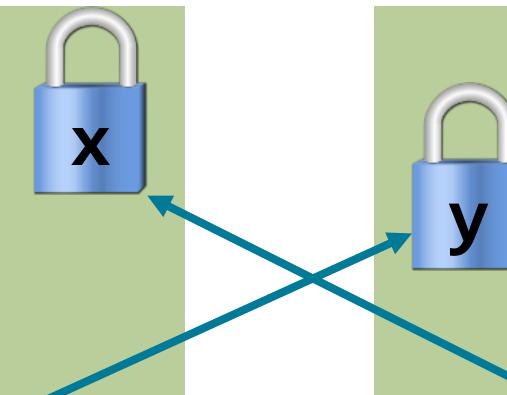
# Deadlocks – Motivation

Suppose **x** and **y** are instances of class **BankAccount**

```
class BankAccount {  
    ...  
    synchronized void withdraw(int amount) {...}  
    synchronized void deposit(int amount) {...}  
    ...  
    synchronized void transferTo(int amount, BankAccount a) {  
        this.withdraw(amount);  
        a.deposit(amount);  
    }  
}
```

Thread 1: `x.transferTo(1,y)`

acquire lock for x  
withdraw from x



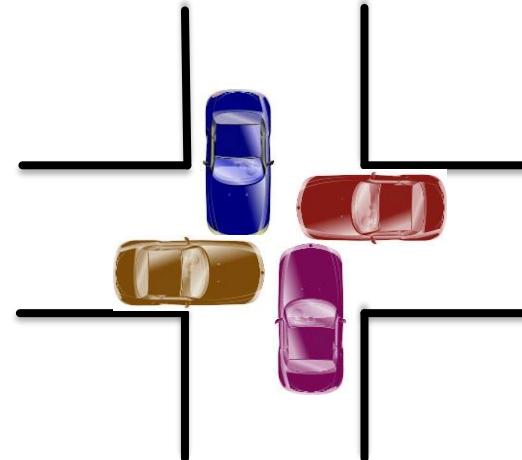
Thread 2: `y.transferTo(1,x)`

acquire lock for y  
withdraw from y  
acquire lock for x

Time

# Deadlocks

**Deadlock:** two or more processes are mutually blocked because each process waits for another of these processes to proceed.



# Threads and Resources

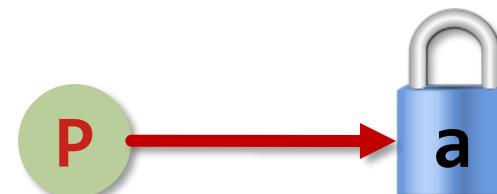
Graphically: Threads



and Resources (Locks)



Thread P *attempts to acquire* resource a:

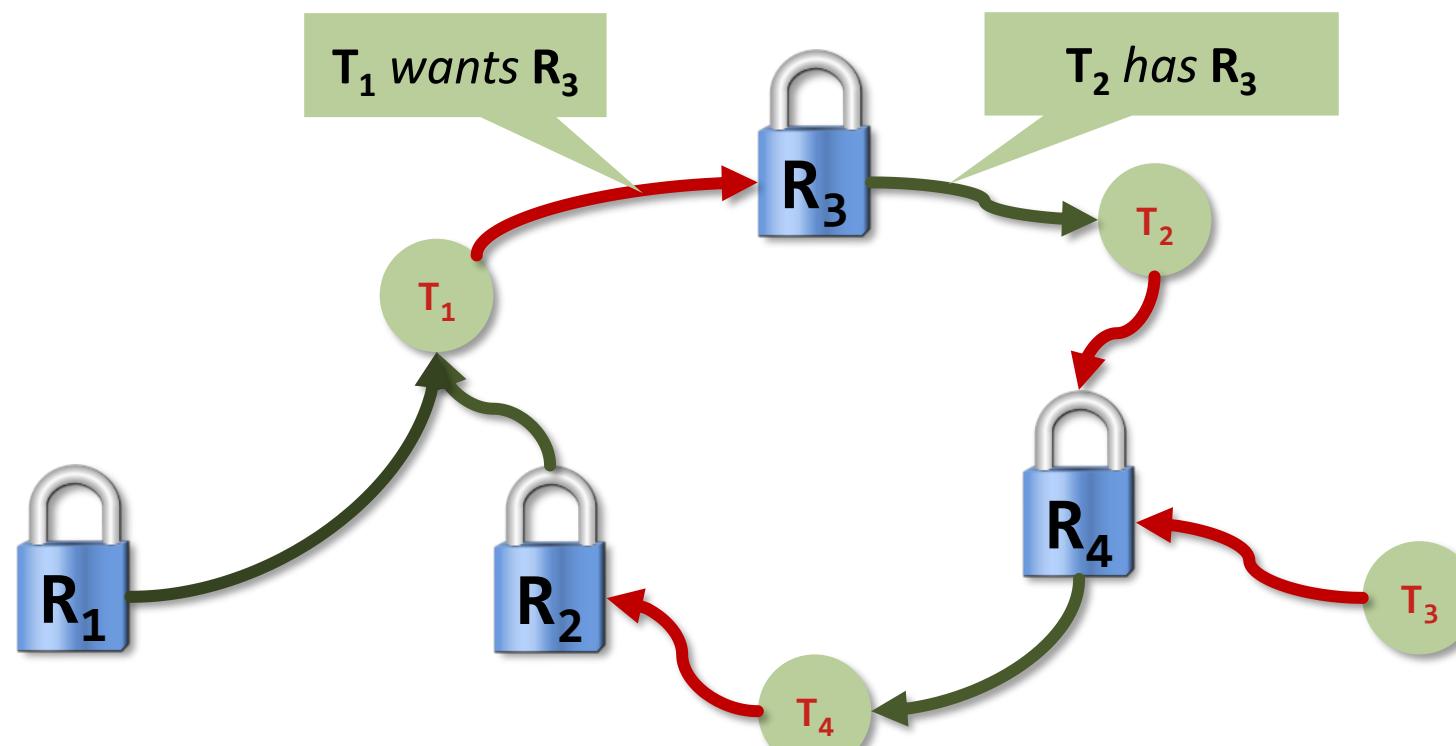


Resource b is *held by* thread Q:



## Deadlocks – more formally

A deadlock for threads  $T_1 \dots T_n$  occurs when the directed graph describing the relation of  $T_1 \dots T_n$  and resources  $R_1 \dots R_m$  contains a cycle.



## Techniques

***Deadlock detection*** in systems is implemented by finding cycles in the dependency graph.

- Deadlocks can, in general, not be healed. Releasing locks generally leads to inconsistent state.

***Deadlock avoidance*** amounts to techniques to ensure a cycle can never arise

- **two-phase locking with retry (release when failed)**
  - Usually in databases where transactions can be aborted without consequence
- **resource ordering**
  - Usually in parallel programming where global state is modified

## Back to our example: what can we do?

```
class BankAccount {  
    ...  
    synchronized void withdraw(int amount) {...}  
    synchronized void deposit(int amount) {...}  
    ...  
    synchronized void transferTo(int amount, BankAccount a) {  
        this.withdraw(amount);  
        a.deposit(amount);  
    }  
}
```

## Option 1: non-overlapping (smaller) critical sections

```
class BankAccount {  
    ...  
    synchronized void withdraw(int amount) {...}  
    synchronized void deposit(int amount) {...}  
    ...  
    void transferTo(int amount, BankAccount a) {  
        this.withdraw(amount);  
        a.deposit(amount);  
    }  
}
```

Money disappears for a (very short?) moment!  
Can we allow such transient inconsistencies?  
**Very often unacceptable!**

## Option 2: one lock for all

```
class BankAccount {  
    static Object globalLock = new Object();  
    // withdraw and deposit protected with globalLock!  
    void withdraw(int amount) {...}  
    void deposit(int amount) {...}  
    ...  
    void transferTo(int amount, BankAccount to) {  
        synchronized (globalLock) {  
            withdraw(amount);  
            to.deposit(amount);  
        }  
    }  
}
```

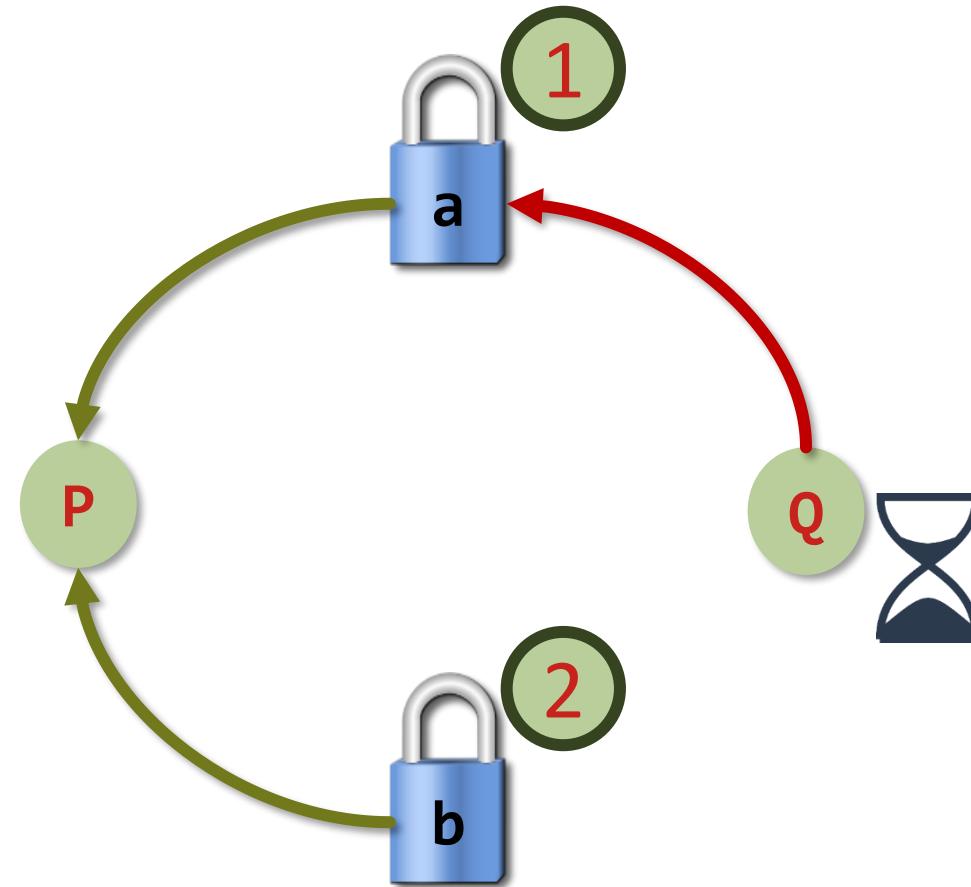
deadlock avoided but no concurrent transfer possible, even not when the pairs of accounts are disjoint.  
Often very inefficient!

## Option 3: global ordering of resources

```
class BankAccount {  
    ...  
    void transferTo(int amount, BankAccount to) {  
        if (to.accountNr < this.accountNr)  
            synchronized(this){  
                synchronized(to) {  
                    withdraw(amount);  
                    to.deposit(amount);  
                }  
            }  
        else  
            synchronized(to){  
                synchronized(this) {  
                    withdraw(amount);  
                    to.deposit(amount);  
                }  
            }  
    }  
}
```

Unique global ordering required.  
Whole program has to obey this order to  
avoid cycles.  
Code taking only one lock can ignore it.

# Ordering of resources



# Programming trick

No globally unique order available? Generate it:

```
class BankAccount {  
    private static final AtomicLong counter = new AtomicLong();  
    private final long index = counter.incrementAndGet();  
    ...  
    void transferTo(int amount, BankAccount to) {  
        if (to.index < this.index)  
            ...  
    }  
}
```

## Another (historic) example: from the Java standard library

```
class StringBuffer {  
    private int count;  
    private char[] value;  
  
    ...  
    synchronized append(StringBuffer sb) {  
        int len = sb.length();  
        if(this.count + len > this.value.length)  
            this.expand(...);  
        sb.getChars(0, len, this.value, this.count);  
    }  
  
    synchronized getChars(int x, int y, char[] a, int z) {  
        "copy this.value[x..y] into a starting at z"  
    }  
}
```

Do you find the two problems?



# Another (historic) example: from the Java standard library

```
class StringBuffer {  
    private int count;  
    private char[] value;  
...  
    synchronized append(StringBuffer sb) {  
        int len = sb.length();  
        if(this.count + len > this.value.length)  
            this.expand(...);  
        sb.getChars(0, len, this.value, this.count);  
    }  
  
    synchronized getChars(int x, int y, char[] a, int z) {  
        "copy this.value[x..y] into a starting at z"  
    }  
}
```

Do you find the two problems?

## Problem #1:

- Lock for **sb** is not held between calls to **sb.length** and **sb.getChars**
- **sb** could get longer
- Would cause **append** to not append whole string
  - The semantics here can be discussed!  
Definitely an issue if **sb** got shorter ☺

## Problem #2:

- Deadlock potential if two threads try to append “crossing” StringBuffers, just like in the bank-account first example
- **x.append(y); y.append(x);**

## Fix?

- **Not easy to fix both problems without extra overheads:**
  - Do not want unique ids on every **StringBuffer**
  - Do not want one lock for all **StringBuffer** objects
- **Actual Java library: initially fixed neither (left code as is; changed javadoc)**
  - Up to clients to avoid such situations with own protocols
- **Today: two classes StringBuffer (claimed to be synchronized) and StringBuilder (not synchronized)**

# Perspective and tricks for programmability

**Code like account-transfer and string-buffer append are difficult to deal with for deadlock**

## 1. Easier case: different types of objects

- Can document a fixed order among types
- Example: “When moving an item from the hashtable to the work queue, never try to acquire the queue lock while holding the hashtable lock”

## 2. Easier case: objects are in an acyclic structure

- Can use the data structure to determine a fixed order
- Example: “If holding a tree node’s lock, do not acquire other tree nodes’ locks unless they are children in the tree”

## Significance of deadlocks

Once understood that (and where) race conditions can occur, with following good programming practice and rules they are relatively easy to cope with.

But the **Deadlock** is **the dominant problem** of reasonably complex concurrent programs or systems and is therefore very important to anticipate!

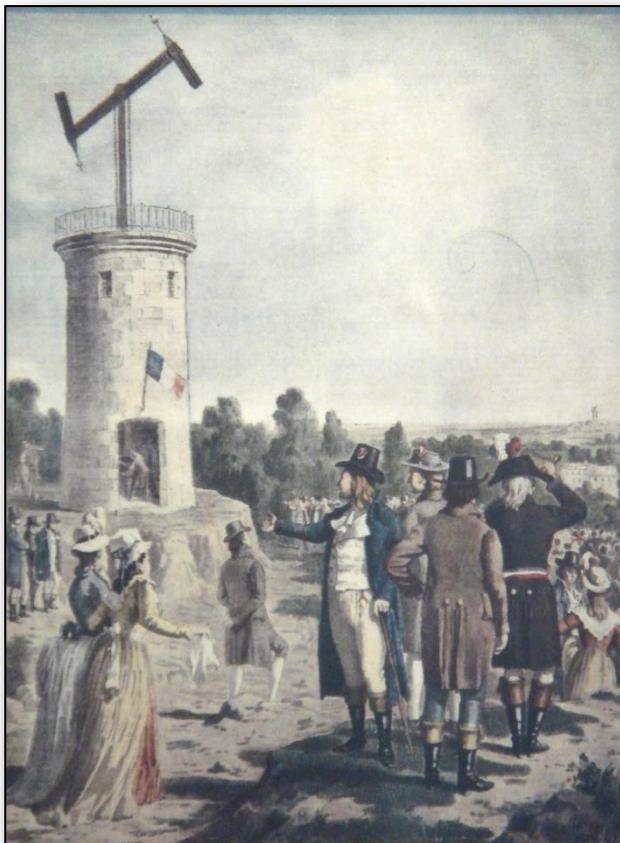
**Starvation** denotes the repeated but unsuccessful attempt of a recently unblocked process to continue its execution.

# Semaphores

## Why do we need more than locks?

- **Locks provide means to enforce atomicity via mutual exclusion**
- **They lack the means for threads to communicate about changes**
  - e.g., changes in the state
- **Thus, they provide no order and are hard to use**
  - e.g., if threads A and B lock object X, it is not determined who comes first
- **Example: producer / consumer queues**

# Semaphore Edsger W. Dijkstra 1965



**Se|maphor**, das od. der; -s, -e [zu griech. σεμα = Zeichen u. φορος = tragend]:  
*Signalmast mit beweglichen Flügeln.*

Optische Telegrafievorrichtung mit Hilfe von schwenkbaren Signalarmen, Claude Chappe 1792

# Semaphore: Semantics

Semaphore: integer-valued abstract data type S with some initial value  $s \geq 0$  and the following operations\*

```
acquire(S)
{
    atomic
        wait until  $S > 0$ 
        dec(S)
}
```

```
release(S)
{
    atomic
        inc(S)
}
```



acquire

(protected)

release

\* Dijkstra called them P (*probeer*), V (*vrijgeven*), also often used: *wait and signal*

## Building a lock with a semaphore

**mutex = Semaphore(1);**

**lock mutex := mutex.acquire()**

only one thread is allowed into the critical section

**unlock mutex := mutex.release()**

one other thread will be let in

**Semaphore number:**

1 → unlocked

0 → locked

$x > 0$  →  $x$  threads will be let into “critical section”

## Example: scaled dot product

When is x ready?

- Execute in parallel:  $x = (\underline{a^T} * \underline{d}) * z$ 
  - a and d are column vectors
  - x, z are scalar
- Assume each vector has 4 elements
  - $x = (a_1 * d_1 + a_2 * d_2 + a_3 * d_3 + a_4 * d_4) * z$
- Parallelize on two processors (using two threads A and B)
  - $x_A = a_1 * d_1 + a_2 * d_2$
  - $x_B = a_3 * d_3 + a_4 * d_4$
  - $x = (x_A + x_B) * z$
- Which synchronization is needed where?
  - Using locks?
  - Using semaphores?

Thread A

```
xA=...;  
lock();  
x=x+xA;  
unlock();
```

Thread B

```
xB=...;  
lock();  
x=x+xB;  
unlock();
```

Thread A

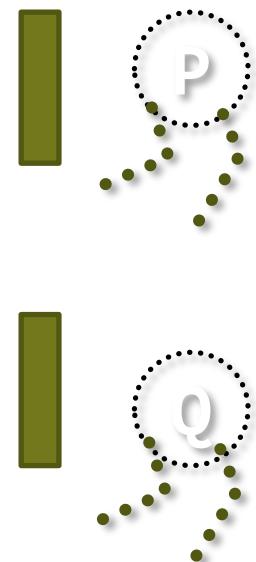
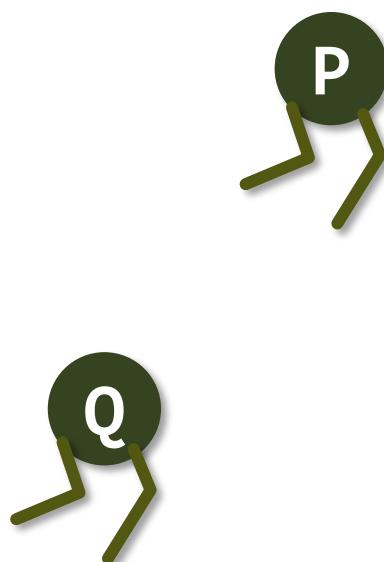
```
xA=...;  
x=x+xA;  
release(S);
```

Thread A

```
xB=...;  
acquire(S);  
x=x+xA;
```

# Rendezvous with Semaphores

- Two processes P and Q executing code.
- Rendezvous: locations in code, where P and Q wait for the other to arrive. Synchronize P and Q.



How would you implement this using Semaphores?

# Rendezvous with Semaphores

Synchronize Processes P and Q at one location (Rendezvous)

Semaphores **P\_Arrived** and **Q\_Arrived**

	P	Q
<i>init</i>	<b>P_Arrived=0</b>	<b>Q_Arrived=0</b>
<i>pre</i>	...	...
<i>rendezvous</i>	?	?
<i>post</i>	...	..

# Rendezvous with Semaphores

Synchronize Processes P and Q at one location (Rendezvous)

Semaphores **P\_Arrived** and **Q\_Arrived**

	P	Q
<i>init</i>	<b>P_Arrived=0</b>	<b>Q_Arrived=0</b>
<i>pre</i>	...	...
<i>rendezvous</i>	<b>release(P_Arrived)</b> ?	<b>acquire(P_Arrived)</b> ?
<i>post</i>	...	...

# Rendezvous with Semaphores

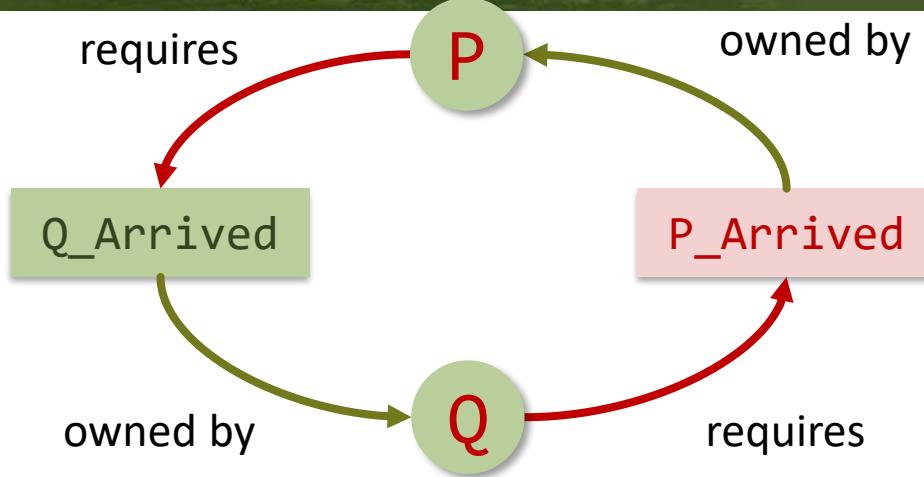
Synchronize Processes P and Q at one location (Rendezvous)

Semaphores **P\_Arrived** and **Q\_Arrived**

Dou you find  
the problem?

	P	Q
<i>init</i>	<b>P_Arrived=0</b>	<b>Q_Arrived=0</b>
<i>pre</i>	...	...
<i>rendezvous</i>	<b>acquire(Q_Arrived)</b> <b>release(P_Arrived)</b>	<b>acquire(P_Arrived)</b> <b>release(Q_Arrived)</b>
<i>post</i>	...	...

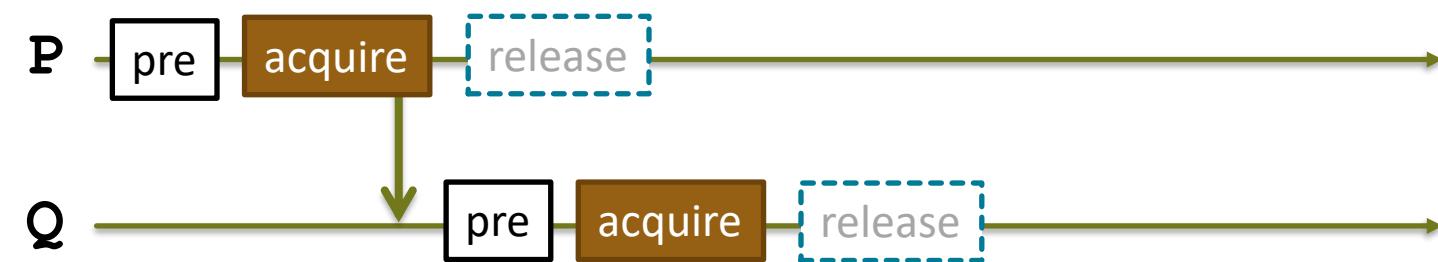
# Deadlock



	P	Q
init	$P\_Arrived=0$	$Q\_Arrived=0$
pre	...	...
rendezvous	$\text{acquire}(Q\_Arrived)$ $\text{release}(P\_Arrived)$	$\text{acquire}(P\_Arrived)$ $\text{release}(Q\_Arrived)$
post	...	...

# Rendezvous with Semaphores

## Wrong solution with Deadlock



# Rendezvous with Semaphores

Synchronize Processes P and Q at one location (Rendezvous)

Assume Semaphores **P\_Arrived** and **Q\_Arrived**

	P	Q
init	<b>P_Arrived=0</b>	<b>Q_Arrived=0</b>
pre	...	...
rendezvous	<b>release(P_Arrived)</b> <b>acquire(Q_Arrived)</b>	<b>acquire(P_Arrived)</b> <b>release(Q_Arrived)</b>
post	...	...

# Digression: Implementing Semaphores without Spinning (blocking queues)

Consider a process list  $Q_s$  associated with semaphore S

**acquire(S)**

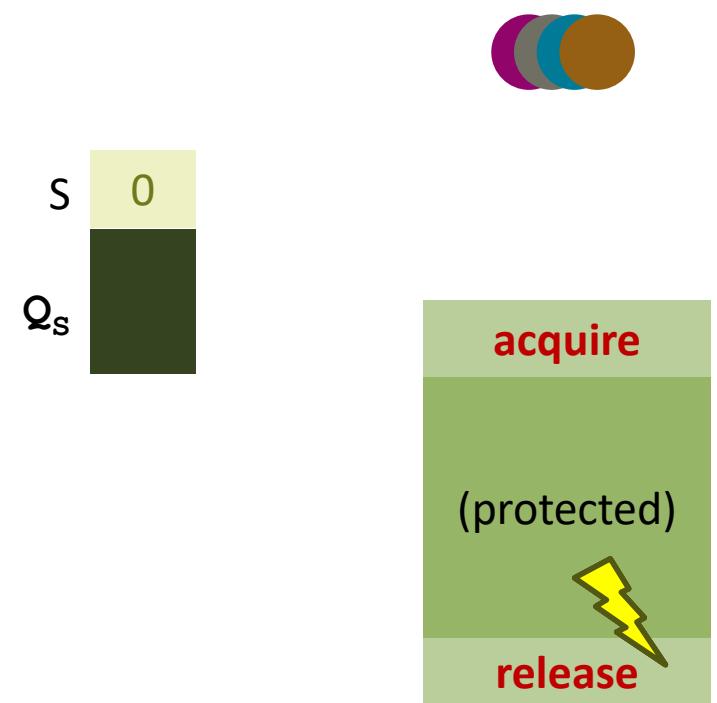
```
{if S > 0 then
    dec(S)
else
    put(Qs, self)
    block(self)
end }
```

atomic

**release(S)**

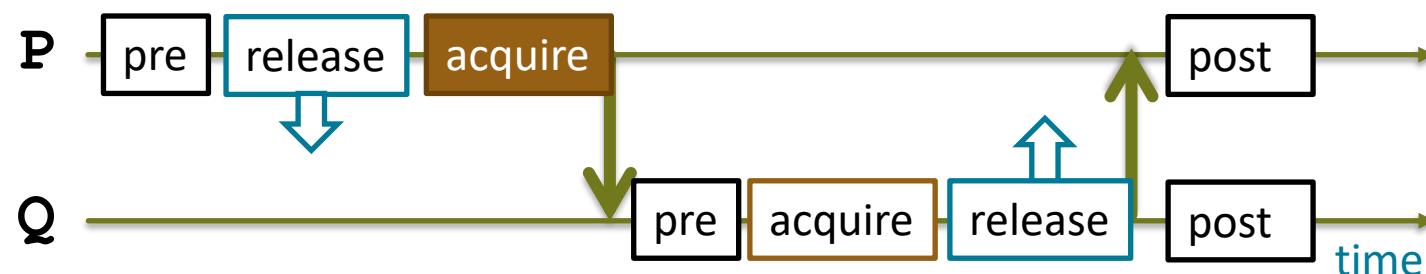
```
{if Qs == Ø then
    inc(S)
else
    get(Qs, p)
    unblock(p)
end }
```

atomic



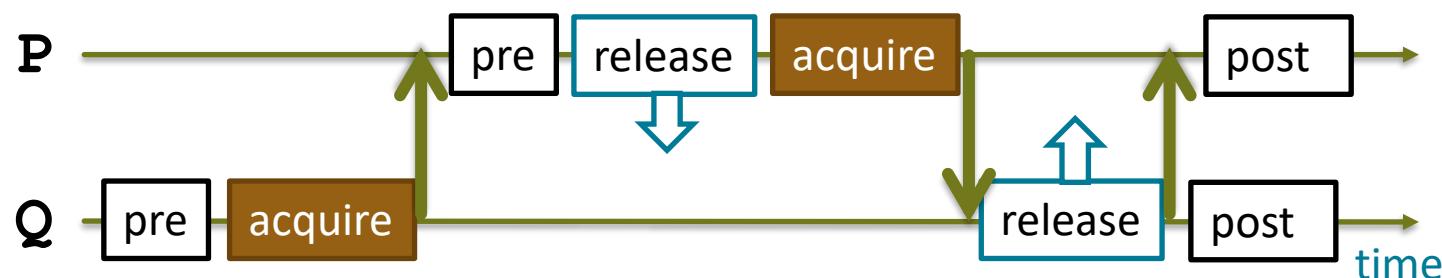
# Scheduling Scenarios

**P first**



release signals (arrow)  
acquire may wait (filled box)

**Q first**



# Rendezvous with Semaphores

Synchronize Processes P and Q at one location (Rendezvous)

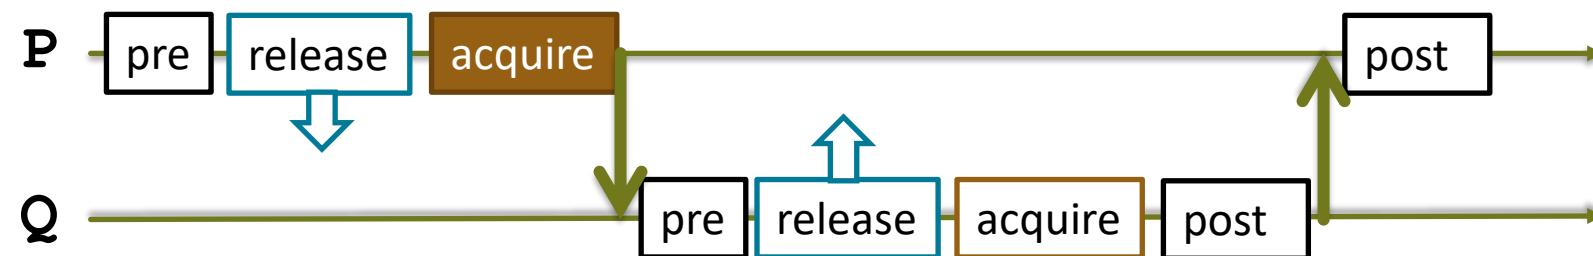
Assume Semaphores **P\_Arrived** and **Q\_Arrived**

	P	Q
init	<code>P_Arrived=0</code>	<code>Q_Arrived=0</code>
pre	<code>...</code>	<code>...</code>
rendezvous	<code>release(P_Arrived)</code> <code>acquire(Q_Arrived)</code>	<code>release(Q_Arrived)</code> <code>acquire(P_Arrived)</code>
post	<code>...</code>	<code>...</code>

# That's even better.

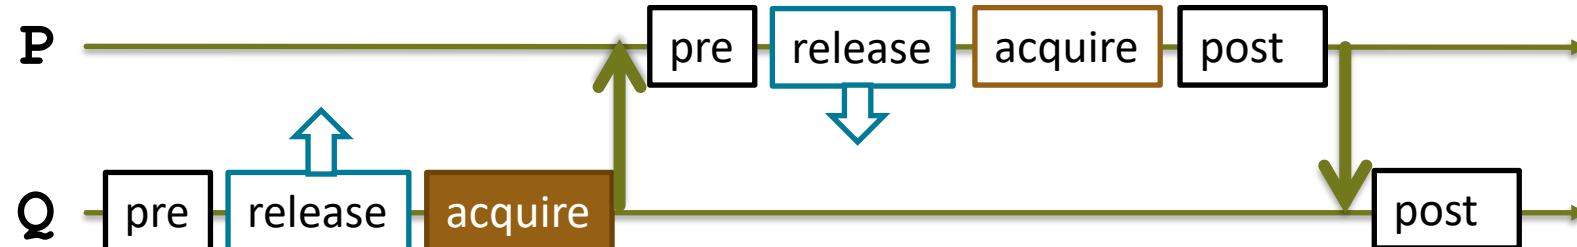
P	Q
<i>init</i>	<code>P_Arrived=0</code>
<i>pre</i>	<code>...</code>
<i>rendezvous</i>	<code>release(P_Arrived)</code> <code>acquire(Q_Arrived)</code>
<i>post</i>	<code>...</code>

## P first



## Q first

release signals (arrow)  
acquire may wait (filled box)



## Back to our dot-product

- **Assume now vectors with 1 million entries on 10,000 threads**
  - Very common! (remember the 57 Pflop/s on 27,360 GPUs on Summit)
  - How would you implement that?
  - Semaphores, locks?
- **Time for a higher-level abstraction!**
  - Supporting threads in bulk-mode  
*Move in lock-step*
  - And enabling a “bulk-synchronous parallel” (BSP) model  
*The full BSP is more complex (supports distributed memory)*



# A Bridging Model for Parallel Computation

The success of the von Neumann model of sequential computation is attributable to the fact that it is an efficient bridge between software and hardware: high-level languages can be efficiently compiled onto this model; yet it can be efficiently implemented in hardware. The author argues that an analogous bridge between software and hardware is required for parallel computation if that is to become as widely used. This article introduces the bulk-synchronous parallel (BSP) model as a candidate for this role, and gives results quantifying its efficiency both in implementing high-level language features and algorithms, as well as in being implemented in hardware.

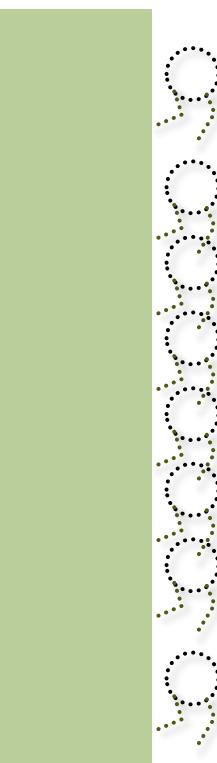
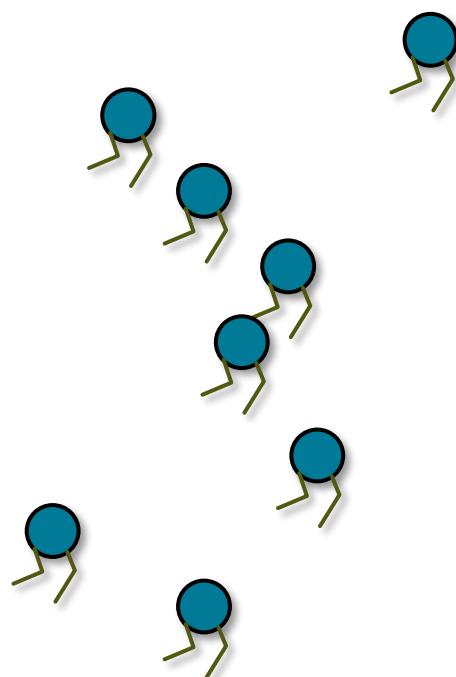
Leslie G. Valiant



# Barriers

# Barrier

Synchronize a number of processes.

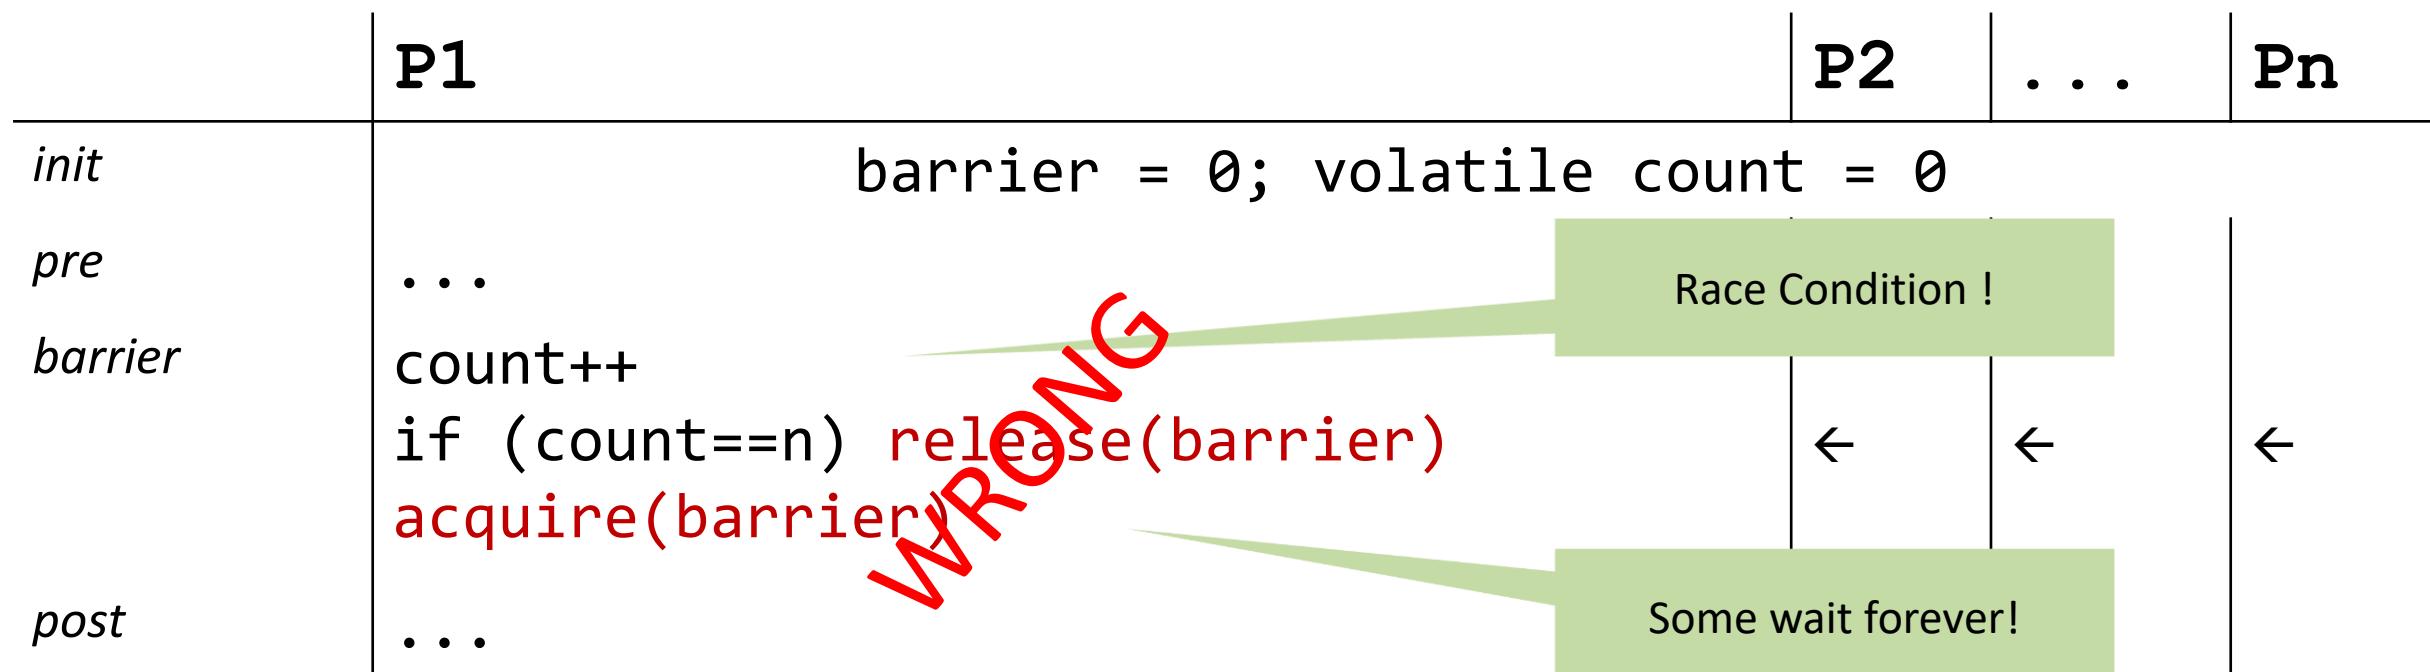


How would you implement this using Semaphores?

## Barrier – 1<sup>st</sup> try

Synchronize a number ( $n$ ) of processes.

Semaphore **barrier**. Integer count .



# Barrier

Synchronize a number ( $n$ ) of processes.

Semaphore **barrier**. Integer count .

P1	
<i>init</i>	barrier = 0; volatile count;
<i>pre</i>	...
<i>barrier</i>	count++ if (count==n) release(barrier) acquire(barrier)
<i>post</i>	...

*WRONG*

Invariants

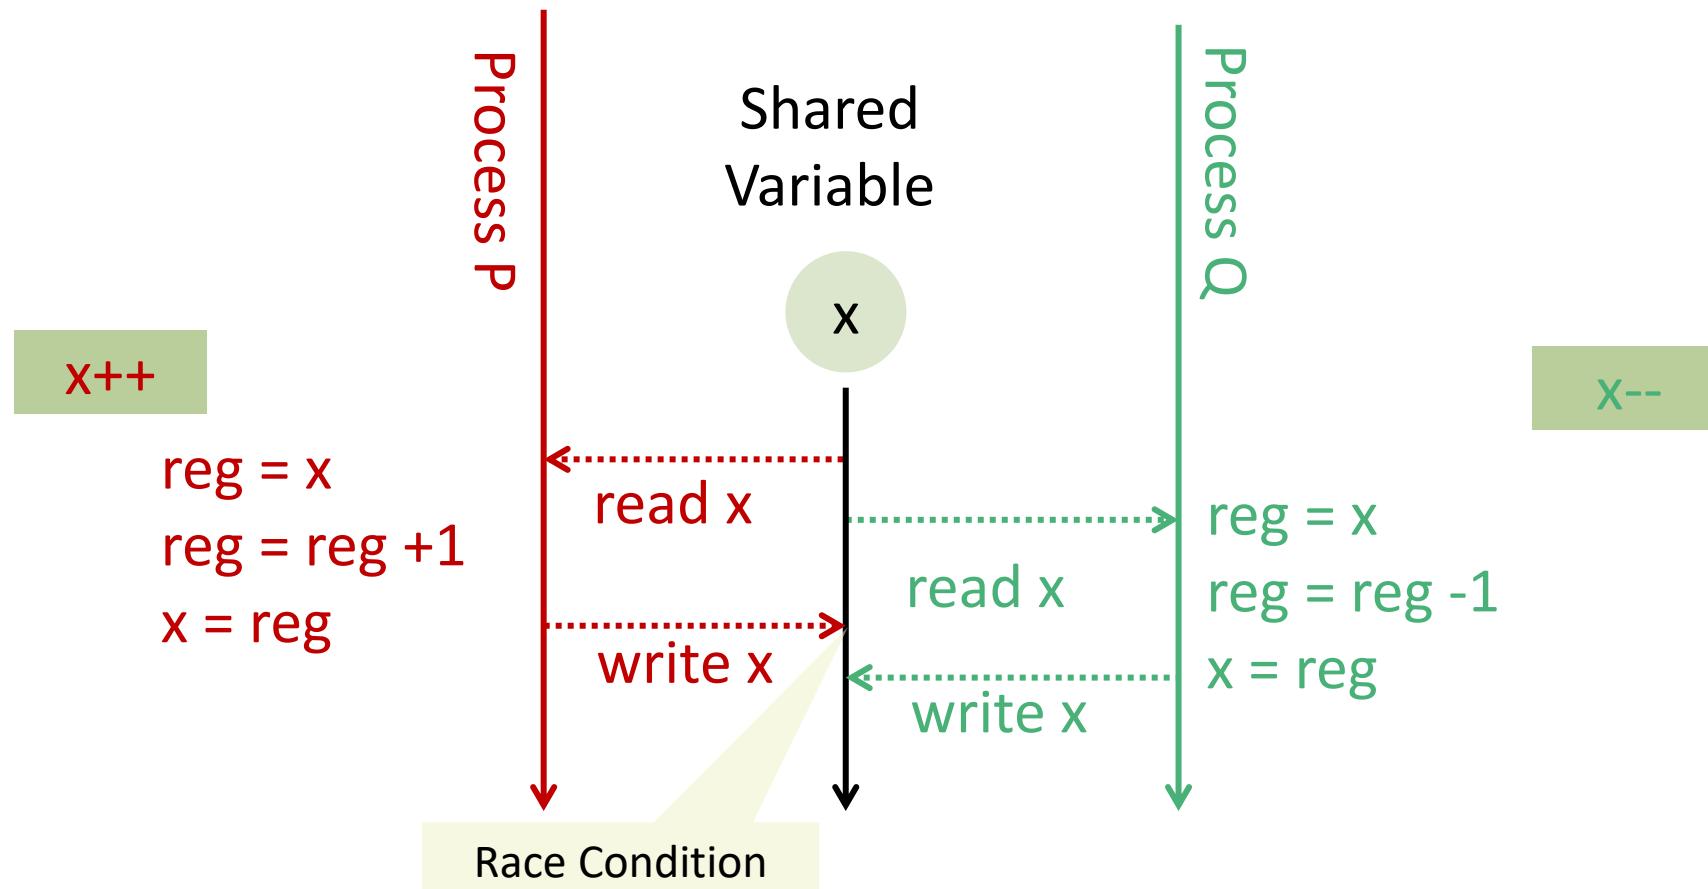
«Each of the processes eventually reaches the acquire statement»

«The barrier will be opened if and only if all processes have reached the barrier»

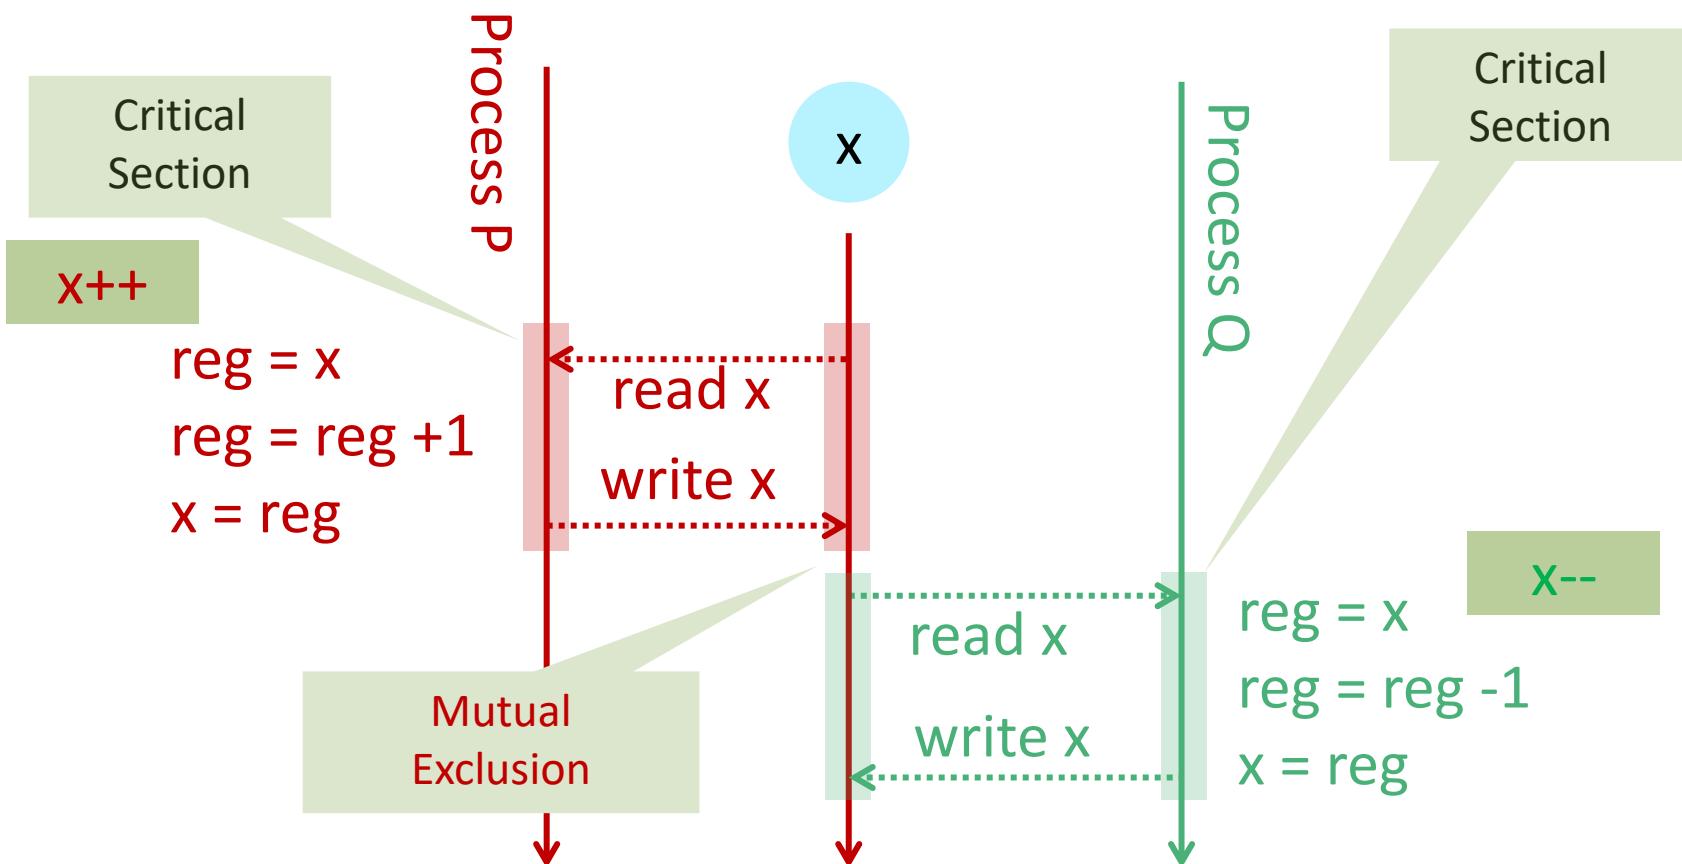
«count provides the number of processes that have passed the barrier» (violated)

«when all processes have reached the barrier then all waiting processes can continue» (violated)

# Recap: Race Condition



# With Mutual Exclusion



# Barrier

Synchronize a number ( $n$ ) of processes.

Semaphores **barrier**, **mutex**. Integer count .

	P1	P2	...	Pn
init	mutex = 1; barrier = 0; count = 0			
pre	...			
barrier	acquire(mutex) count++ release(mutex) if (count==n) release(barrier) acquire(barrier) release(barrier)			
post	...	←	←	←



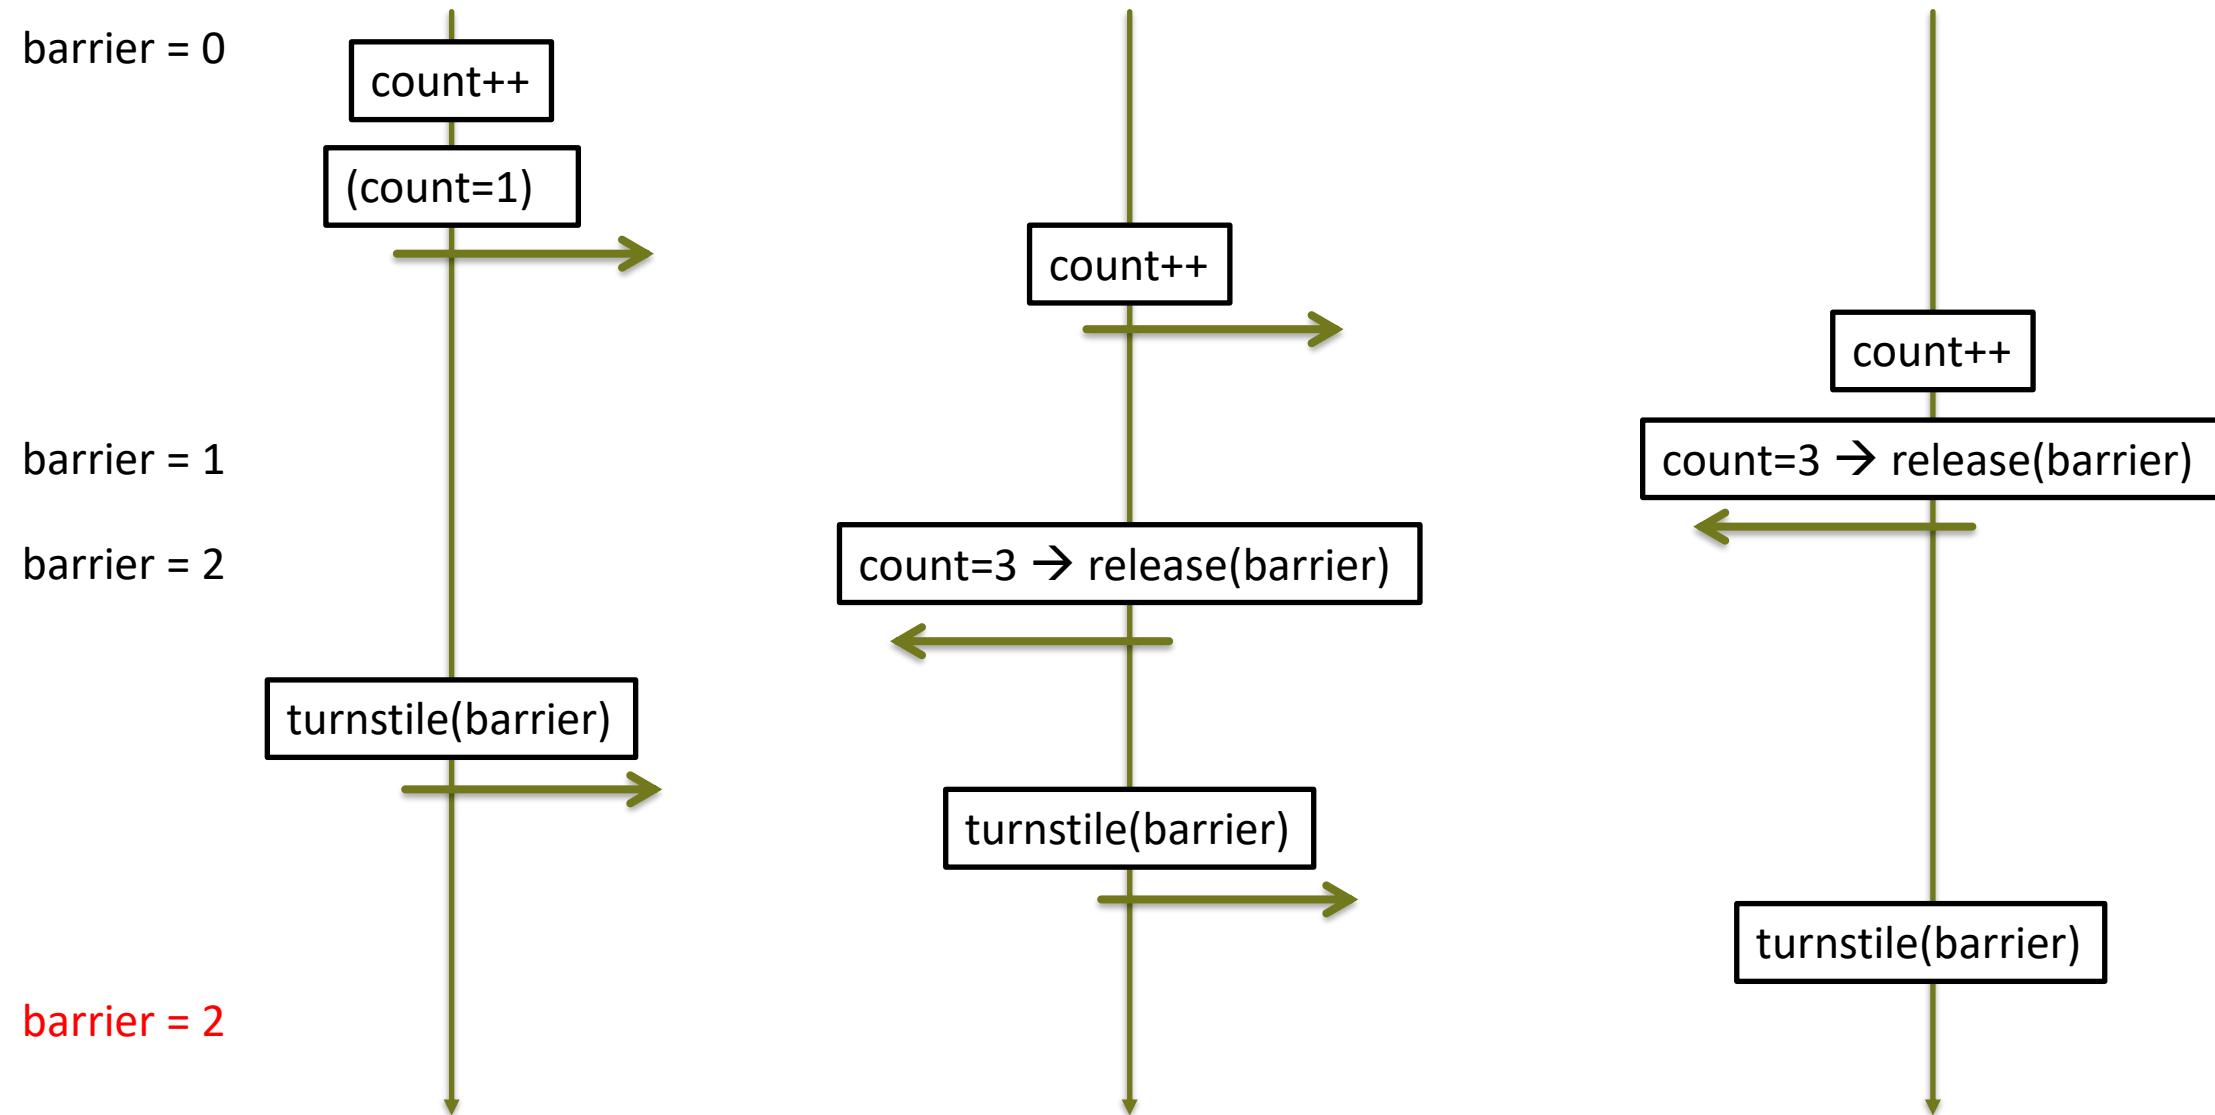
# Reusable Barrier. 1st trial.

	P1	...	Pn
init		mutex = 1; barrier = 0; count = 0	
pre	...		
barrier	acquire(mutex) count++ release(mutex) if (count==n) release(barrier)		Dou you see the problem?
	acquire(barrier) release(barrier)	Race Condition !	←
	acquire(mutex) count-- release(mutex) if (count==0) acquire(barrier)	Race Condition !	←
post	...		

# Reusable Barrier. 1st trial.

	P1	...	Pn
init		mutex = 1; barrier = 0; count = 0	
pre	...		
barrier	acquire(mutex) count++ release(mutex) if (count==n) release(barrier)  acquire(barrier) release(barrier)  acquire(mutex) count-- release(mutex) if (count==0) acquire(barrier)  ...	Invariants  «Only when all processes have reached the turnstyle it will be opened the first time»  «When all processes have run through the barrier then count = 0»  «When all processes have run through the barrier then barrier = 0» (violated)	←
post			

# Illustration of the problem: scheduling scenario



# Reusable Barrier. 2nd trial.

	P1	...	Pn
init		mutex = 1; barrier = 0; count = 0	
pre	...		
barrier	acquire(mutex) count++ if (count==n) release(barrier) release(mutex)		
	acquire(barrier) release(barrier)	Process can pass other processes!	
	acquire(mutex) count-- if (count==0) acquire(barrier) release(mutex)		<
post	...		<

Do you see the problem?

Process can pass other processes!

# Reusable Barrier. 2nd trial.

	P1	...	Pn
init		mutex = 1; barrier = 0; count = 0	
pre	...		
barrier	acquire(mutex) count++ if (count==n) release(barrier) release(mutex)		
	acquire(barrier) release(barrier)		
	acquire(mutex) count-- if (count==0) acquire(barrier) release(mutex)		
post	...		



## Invariants

«When all processes have passed the barrier, it holds that barrier = 0»

« Even when a single process has passed the barrier, it holds that barrier = 0 » (violated)

# Solution: Two-Phase Barrier

```
init      mutex=1; barrier1=0; barrier2=1; count=0

barrier
    acquire(mutex)
        count++;
        if (count==n)
            acquire(barrier2); release(barrier1)
    release(mutex)

    acquire(barrier1); release(barrier1);
    // barrier1 = 1 for all processes, barrier2 = 0 for all processes
    acquire(mutex)
        count--;
        if (count==0)
            acquire(barrier1); release(barrier2)
    signal(mutex)

    acquire(barrier2); release(barrier2)
    // barrier2 = 1 for all processes, barrier1 = 0 for all processes
```

## Lesson Learned ?

- **Semaphore, Rendevouz and Barrier:**
- **Concurrent programming is prone to errors in reasoning.**
- **A naive approach with trial and error is close-to impossible.**
- **Ways out:**
  - Identify **invariants** in the problem domain, ensure they hold for your implementation
  - Identify and apply **established patterns**
  - Use known **good libraries** (like in the Java API)

## Summary

**Locks are not enough: we need methods to wait for events / notifications**

**Semaphores**

**Rendezvous and Barriers**

**Next lecture:**

**Producer-Consumer Problem**

**Monitors and condition variables**